

Google Classroom Code: mhxgl24

PyTorch

Deep Learning (DS-5006)

Dr. Adeel Mumtaz

Lecture 4

Fall, 2022




National University
Of Computer and Emerging Sciences

Contents

- Basics of NN Review
 - A Simple Regression Problem
 - Training Loop
 - NN Summary
 - Home Task
- **Quiz-1**
- Multiple Regression
 - Vectorization of Gradient Descent
- PyTorch
 - Introduction
 - History
 - Features
 - Tensorflow vs Pytorch
 - Tensors
 - Creating
 - Reshaping
- Copying
 - Numpy conversion
 - GPU Tensors
 - Creating Parameters
 - Dynamic Computation Graph
 - Autograd
 - Optimizers
 - Loss Functions
 - Model Class
 - Nested Models
- A Simple Regression Problem (Pytorch Implementation)
- Summary
- Graded Home Task (backprop)

Simple Linear Regression



The Linear Unit: $y = wx + b$

- **NN=Architecture + Parameters**
- **Training NN = Given data learn best Parameters which gives minimum loss (usually an iterative process/algorithm)**

Simple Linear Regression

$$\mathcal{L}(x, b, w) = \frac{1}{N} \sum_{i=1}^N (wx_i + b - y_i)^2$$

$$dw = \frac{\partial \mathcal{L}(x, b, w)}{\partial w} = 2 * \frac{1}{N} \sum_{i=1}^N (wx_i + b - y_i)(x_i)$$

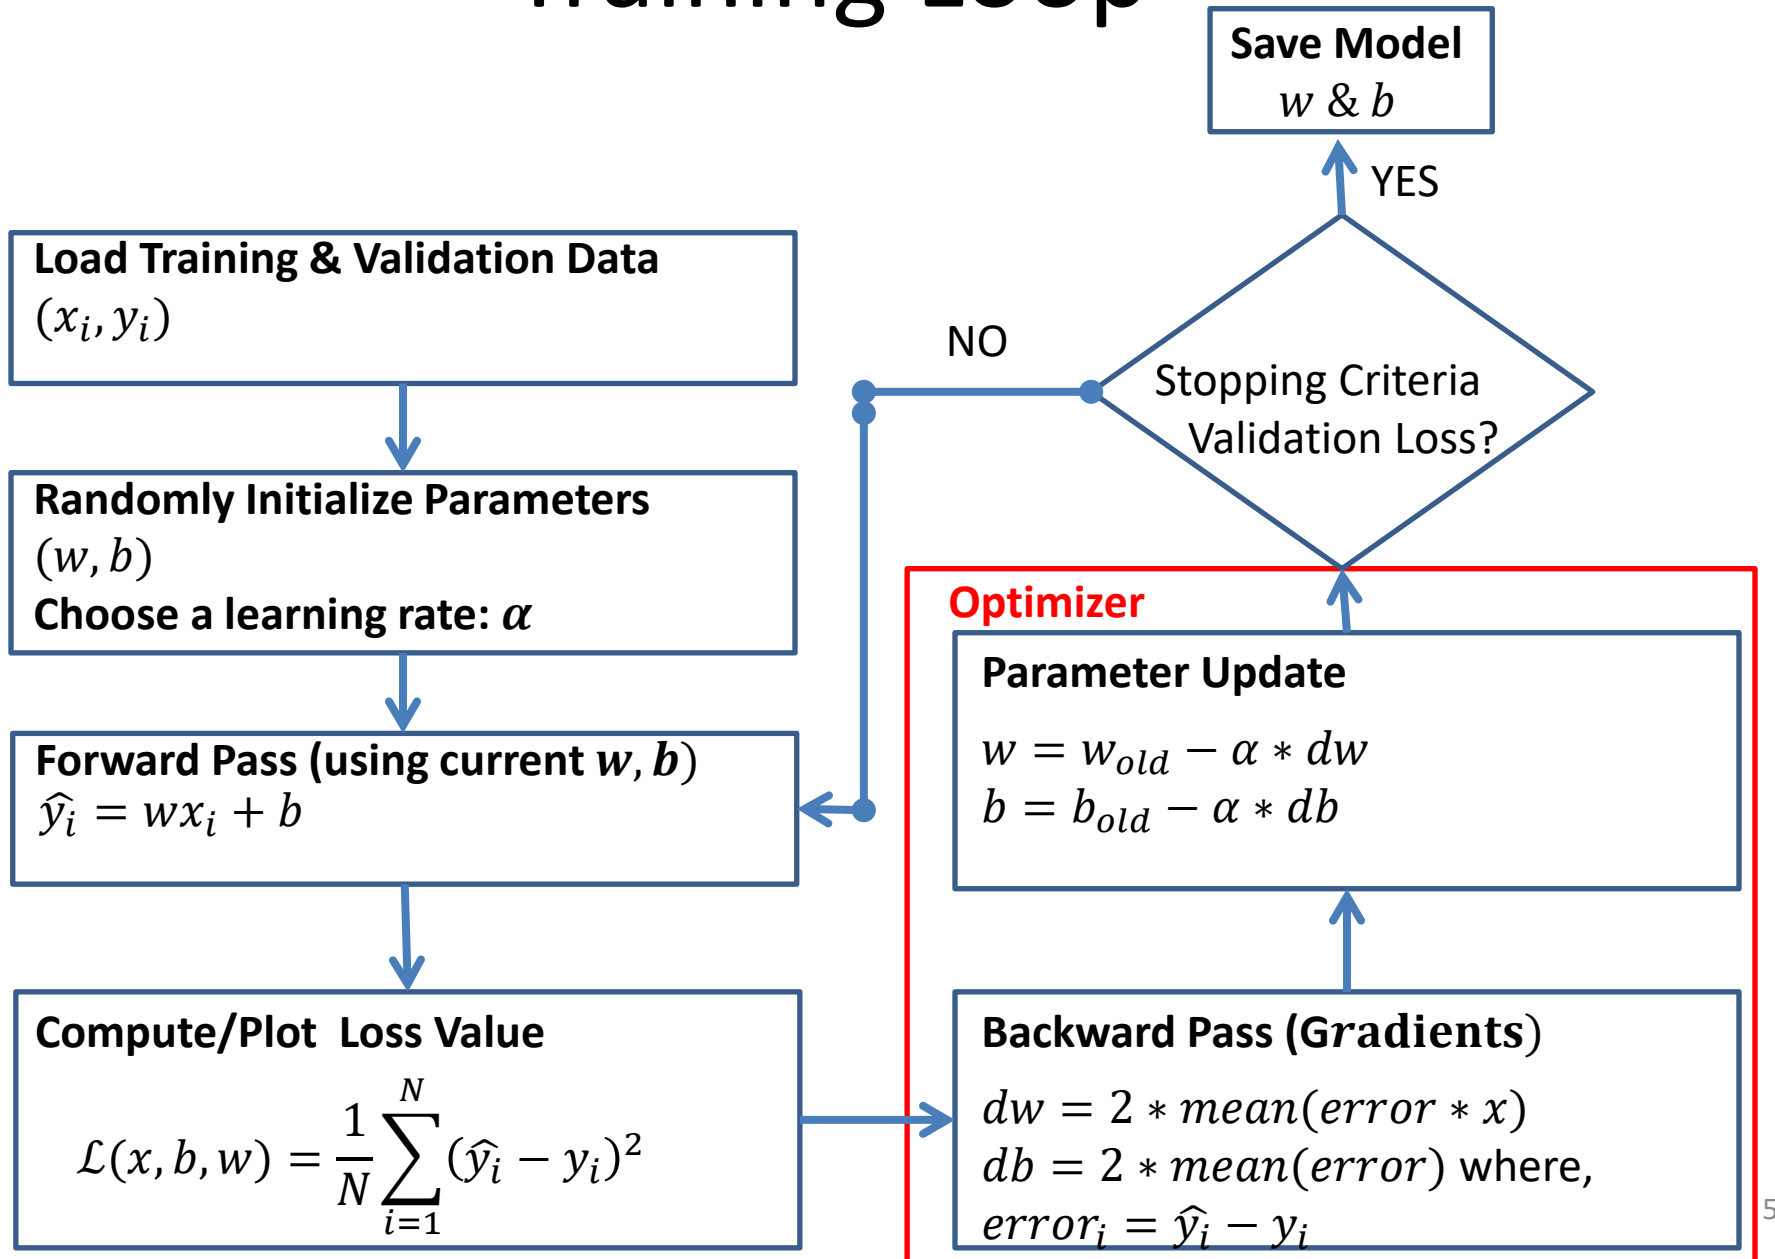
$$= 2 * \frac{1}{N} \sum_{i=1}^N (error_i)(x_i) = 2 * mean(error * x)$$

$$db = \frac{\partial \mathcal{L}(x, b, w)}{\partial b} = 2 * \frac{1}{N} \sum_{i=1}^N (error_i) = 2 * mean(error)$$

$$w_{new} = w_{old} - \alpha * dw$$

$$b_{new} = b_{old} - \alpha * db$$

Training Loop



Training Loop

```
39 #training loop
40 #initializing parameters
41 trainLosses=[]
42 valLosses=[]
43 lr=0.1
44 w=np.random.randn(1)
45 b=np.random.randn(1)
46 for i in range(100):
47     #forward pass
48     yhat=w*x_train+b #note vectorized operation
49     #MSE loss
50     error=yhat-y_train
51     loss= (error**2).mean()
52     trainLosses.append(loss)
53     #computing gradients
54     db=2*error.mean()
55     dw=2*(x_train*error).mean()
56     #weight update
57     b=b-lr*db
58     w=w-lr*dw
```

NN Summary Review

- Data Set, Training, Validation, Test
- NN as function Approximators
- Cost/Loss Function
 - MSE Loss for regression
- Architecture of NN
- Parameters
- Training Loop
- Optimizer
- Learning Rate
- Types of Gradient Descent
 - Epoch
 - Batch
- Loading/Saving Model

Home Task

- Compare learning curves for different values of learning rate
- Convert code from Batch Gradient Descent to Stochastic Gradient Descent and compare learning curves

Quiz-1 (10min)

- Consider a regression problem with one input variable x and one output variable y .
- Following regression model/architecture having unknown parameter w is assumed:
- $y_i = \log(w^2 x_i)$
- Derive the **batch gradient descent update equation** for parameter w assuming MSE Loss

Solution

$$\mathcal{L}(x, b, w) = \frac{1}{N} \sum_{i=1}^N (\log(w^2 x_i) - y_i)^2$$

$$dw = \frac{\partial \mathcal{L}(x, b, w)}{\partial w} = 2 * \frac{1}{N} \sum_{i=1}^N (\log(w^2 x_i) - y_i) * \frac{1}{w^2 x_i} * 2w x_i$$

$$dw = 4 * \frac{1}{N} \sum_{i=1}^N (\log(w^2 x_i) - y_i) * \frac{1}{w}$$

$$dw = \frac{4}{N * w} \sum_{i=1}^N (\log(w^2 x_i) - y_i)$$

$$w_{new} = w_{old} - \alpha * dw$$


MULTIPLE REGRESSION

Multiple regression

- Multiple regression is a technique that can be used to analyze the relationship between a single dependent variable and several independent variables

Size (feet ²) x_1	Number of bedrooms x_2	Number of floors x_3	Age of home (years) x_4	Price (\$1000) y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

Model



Vectorization

$$X = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \cdots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \cdots & x_n^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} & \cdots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \cdots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} & \cdots & x_n^{(m)} \end{bmatrix}_{m \times (n+1)}$$

$$\hat{Y} = XW$$

$$W = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_n \end{bmatrix}_{n+1}$$

$$Y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}_m$$

Gradient Descent (Vectorization)

$$\mathcal{L}(X, W) = \frac{1}{m} \sum_{i=1}^m (W^T x^{(i)} - y^{(i)})^2$$

$$dW = \frac{\partial \mathcal{L}(X, W)}{\partial W} = \frac{2}{m} \sum_{i=1}^m (x^{(i)} W - y^{(i)}) x^{(i)}$$

$$W_{new} = W_{old} - \alpha * dW$$

$$b_{new} = ?$$

Scalar derivative			Vector derivative		
$f(x)$	\rightarrow	$\frac{df}{dx}$	$f(\mathbf{x})$	\rightarrow	$\frac{df}{d\mathbf{x}}$
bx	\rightarrow	b	$\mathbf{x}^T \mathbf{B}$	\rightarrow	\mathbf{B}
bx	\rightarrow	b	$\mathbf{x}^T \mathbf{b}$	\rightarrow	\mathbf{b}
x^2	\rightarrow	$2x$	$\mathbf{x}^T \mathbf{x}$	\rightarrow	$2\mathbf{x}$
bx^2	\rightarrow	$2bx$	$\mathbf{x}^T \mathbf{B} \mathbf{x}$	\rightarrow	$2\mathbf{B} \mathbf{x}$

Gradient Descent (Matrixization)

$$\mathcal{L}(X, W) = \frac{1}{m} \sum_{i=1}^m (W^T x^{(i)} - y^{(i)})^2$$

$$dW = \frac{\partial \mathcal{L}(X, W)}{\partial W} = \frac{2}{m} X^T (XW - Y)$$

$$W_{new} = W_{old} - \alpha * dW$$

$$\mathcal{L}(X, W) = \frac{1}{m} (XW - Y)^T (XW - Y)$$

$$\mathcal{L}(X, W) = \frac{1}{m} ((XW)^T - Y^T)(XW - Y)$$

$$\mathcal{L}(X, W) = \frac{1}{m} (W^T X^T - Y^T)(XW - Y)$$

$$\mathcal{L}(X, W) = \frac{1}{m} (W^T X^T XW - W^T X^T Y - Y^T XW + Y^T Y)$$

$$dW = \frac{\partial \mathcal{L}(X, W)}{\partial W} = \frac{1}{m} (2X^T XW - X^T Y - X^T Y)$$

$$dW = \frac{\partial \mathcal{L}(X, W)}{\partial W} = \frac{2}{m} (X^T XW - X^T Y)$$

Scalar derivative		Vector derivative	
$f(x)$	$\rightarrow \frac{df}{dx}$	$f(\mathbf{x})$	$\rightarrow \frac{df}{d\mathbf{x}}$
bx	$\rightarrow b$	$\mathbf{x}^T \mathbf{B}$	$\rightarrow \mathbf{B}$
bx	$\rightarrow b$	$\mathbf{x}^T \mathbf{b}$	$\rightarrow \mathbf{b}$
x^2	$\rightarrow 2x$	$\mathbf{x}^T \mathbf{x}$	$\rightarrow 2\mathbf{x}$
bx^2	$\rightarrow 2bx$	$\mathbf{x}^T \mathbf{B} \mathbf{x}$	$\rightarrow 2\mathbf{B} \mathbf{x}$

TIME TO TORCH IT :-)



- <https://pytorch.org/>
- PyTorch is a deep learning framework and scientific computing package based on Python that uses the power of graphics processing units (GPU)
- FROM RESEARCH TO PRODUCTION
 - An open source machine learning framework that accelerates the path from research prototyping to production deployment



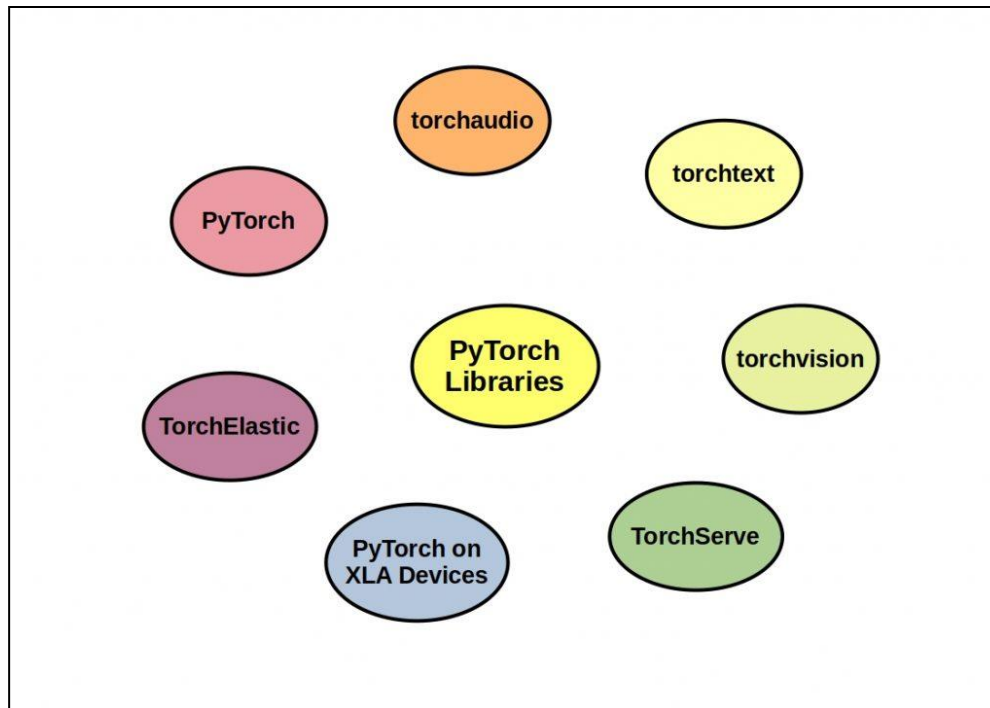
- Primarily developed by Facebook's AI Research lab (FAIR)
- PyTorch was launched in October of 2016 as Torch
- Caffe2 was merged into PyTorch at the end of March 2018
- A number of pieces of Deep Learning software are built on top of PyTorch, including:
 - Tesla Autopilot.
 - Uber's Pyro.
 - HuggingFace's Transformers.
 - PyTorch Lightning.
 - Catalyst.



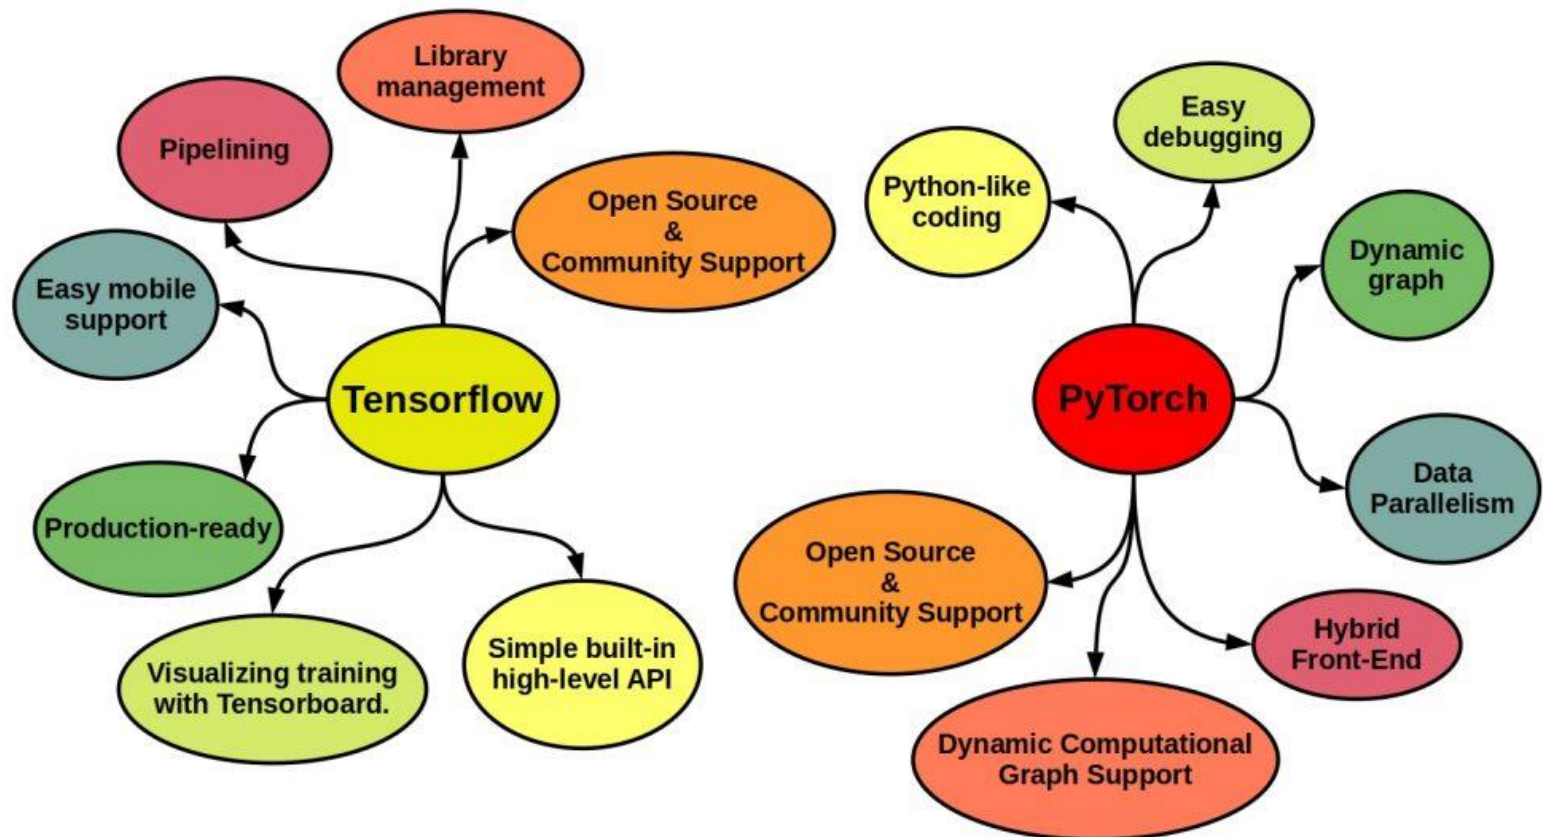
- Distributed Training
- ONNX (Open Neural Network Exchange) Support
- Autograd module
- Optim module
 - Most of the commonly used methods are already supported,
- nn module
 - layers and tools to easily create a neural networks by just defining the layers of the network.




- **PyTorch rules research**
- It has become the framework of choice at [CVPR](#), [ICLR](#), and [ICML](#), among others



PyTorch vs Tensorflow



PyTorch vs Tensorflow



TENSORS


Tensor

- Scalar (a single number) has zero dimensions
- Vector has one dimension
- Matrix has two dimensions
- Tensor has three or more dimensions
- But, to keep things simple, it is commonplace to call vectors and matrices tensors as well
- so, from now on, **everything is either a scalar or a tensor**

Tensor




Scalar



Vector



Matrix



Tensor

Scalar



Vector



Matrix



Tensor



Creating Tensors

```
import torch
```

```
scalar = torch.tensor(3.14159)
vector = torch.tensor([1, 2, 3])
matrix = torch.ones((2, 3), dtype=torch.float)
tensor = torch.randn((2, 3, 4), dtype=torch.float)

print(scalar)
print(vector)
print(matrix)
print(tensor)
```

Creating Tensors

```
tensor(3.1416)
tensor([1, 2, 3])
tensor([[1., 1., 1.],
        [1., 1., 1.]])
tensor([[-1.0658, -0.5675, -1.2903, -0.1136],
        [ 1.0344,  2.1910,  0.7926, -0.7065],
        [ 0.4552, -0.6728,  1.8786, -0.3248]],

        [[-0.7738,  1.3831,  1.4861, -0.7254],
         [ 0.1989, -1.0139,  1.5881, -1.2295],
         [-0.5338, -0.5548,  1.5385, -1.2971]]])
```

Getting Shape of Tensor

```
print(tensor.size(), tensor.shape)
```

```
torch.Size([2, 3, 4]) torch.Size([2, 3, 4])
```

Reshaping a Tensor

- You can also reshape a tensor using its **view()** (preferred) or **reshape()** methods
- Beware it **DOES NOT** create a new, independent, tensor!

```
# We get a tensor with a different shape but it still is  
# the SAME tensor  
same_matrix = matrix.view(1, 6)  
# If we change one of its elements...  
same_matrix[0, 1] = 2.  
# It changes both variables: matrix and same_matrix  
print(matrix)  
print(same_matrix)
```

```
tensor([[1., 2., 1.],  
        [1., 1., 1.]])  
tensor([[1., 2., 1., 1., 1., 1.]])
```

Copying a Tensor

- If you want to copy all data for real, that is, duplicate the data in memory, you may
- use either its `new_tensor()` or `clone()` methods

```
# Lets follow PyTorch's suggestion and use "clone" method
another_matrix = matrix.view(1, 6).clone().detach()
# Again, if we change one of its elements...
another_matrix[0, 1] = 4.
# The original tensor (matrix) is left untouched!
print(matrix)
print(another_matrix)
```

```
tensor([[1., 2., 1.],
        [1., 1., 1.]])
tensor([[1., 4., 1., 1., 1., 1.]])
```

Numpy to PyTorch

- `as_tensor()` or `from_numpy()` is used to convert Numpy arrays to PyTorch Tensors
- This operation preserves the type of the array
- both `as_tensor()` and `from_numpy()` return a tensor that **shares the underlying data with the original Numpy array**.
- `torch.tensor()` always **makes a copy of Numpy array**

```
x_train_tensor = torch.as_tensor(x_train)
x_train.dtype, x_train_tensor.dtype
```

```
(dtype('float64'), torch.float64)
```

```
dummy_array = np.array([1, 2, 3])
dummy_tensor = torch.as_tensor(dummy_array)
# Modifies the numpy array
dummy_array[1] = 0
# Tensor gets modified too...
dummy_tensor
```

```
tensor([1, 0, 3])
```


PyTorch to Numpy

You can also perform the opposite operation, namely, transforming a PyTorch tensor back to a *Numpy* array. That's what `numpy()` is good for:

```
dummy_tensor.numpy()
```

```
array([1, 0, 3])
```

GPU Tensors

- These tensors store their data in the graphics card's memory
- Operations on top of them are performed by the GPU
- If you have a graphics card from **NVIDIA**, you can use the power of its GPU to speed up model training.
- PyTorch supports the use of these GPUs for model training using CUDA (Compute Unified Device Architecture), which needs to be previously installed and configured

PyTorch Build	Stable (1.12.1)		Preview (Nightly)		LTS (1.8.2)
Your OS	Linux		Mac		Windows
Package	Conda		Pip	LibTorch	Source
Language	Python			C++ / Java	
Compute Platform	CUDA 10.2	CUDA 11.3	CUDA 11.6	ROCm 5.1.1	CPU
Run this Command:	NOTE: 'conda-forge' channel is required for cudatoolkit 11.6 conda install pytorch torchvision torchaudio cudatoolkit=11.6 -c pytorch -c conda-forge				

GPU Tensors

- you should always make your code GPU-ready,
- that is, it should automatically run in a GPU, if **one is available**

```
device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

"Why cuda:0? Are there others, like cuda:1, cuda:2 and so on?"

```
n_cudas = torch.cuda.device_count()  
for i in range(n_cudas):  
    print(torch.cuda.get_device_name(i))
```

GeForce GTX 1060 6GB

GPU Tensors

- Moving tensor to GPU using `.to()`

```
gpu_tensor = torch.as_tensor(x_train).to(device)  
gpu_tensor[0]
```

```
tensor([0.7713], device='cuda:0', dtype=torch.float64)
```

"Should I use `to(device)`, even if I am using CPU only?"

Putting it all together

```
device = 'cuda' if torch.cuda.is_available() else 'cpu'  
  
# Our data was in Numpy arrays, but we need to transform them  
# into PyTorch's Tensors and then we send them to the  
# chosen device  
x_train_tensor = torch.as_tensor(x_train).float().to(device)  
y_train_tensor = torch.as_tensor(y_train).float().to(device)
```

```
# Here we can see the difference - notice that .type() is more  
# useful since it also tells us WHERE the tensor is (device)  
print(type(x_train), type(x_train_tensor), x_train_tensor.type())
```

```
<class 'numpy.ndarray'> <class 'torch.Tensor'>  
torch.cuda.FloatTensor
```

GPU Tensors to Numpy

```
back_to_numpy = x_train_tensor.numpy()
```

TypeError: can't convert CUDA tensor to numpy. Use Tensor.cpu() to copy the tensor to host memory first.

```
back_to_numpy = x_train_tensor.cpu().numpy()
```

Summary So Far

```
scalar = torch.tensor(3.14159)
vector = torch.tensor([1, 2, 3])
matrix = torch.ones((2, 3), dtype=torch.float)
tensor = torch.randn((2, 3, 4), dtype=torch.float)
```

```
print(tensor.size(), tensor.shape)
```

```
same_matrix = matrix.view(1, 6)
```

```
another_matrix = matrix.view(1, 6).clone().detach()
```

```
x_train_tensor = torch.as_tensor(x_train)
```

```
dummy_tensor.numpy()
```

```
device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
x_train_tensor = torch.as_tensor(x_train).float().to(device)
```

```
back_to_numpy = x_train_tensor.cpu().numpy()
```

Creating Parameters

- What is difference between a tensor used for training (validation or test) data and a tensor used for the parameter/weight of NN?
- Parameters requires computation of gradients with respect to a loss function during training
- A tensor for a learnable parameter requires a gradient!
- Always assign tensors to a device at the moment of their creation to avoid unexpected behaviors!

Creating Parameters (best approach)

```
# FINAL
# We can specify the device at the moment of creation
# RECOMMENDED!

# Step 0 - Initializes parameters "b" and "w" randomly
torch.manual_seed(42)
b = torch.randn(1, requires_grad=True, \
                dtype=torch.float, device=device)
w = torch.randn(1, requires_grad=True, \
                dtype=torch.float, device=device)
print(b, w)
```


```
tensor([0.1940], device='cuda:0', requires_grad=True)
tensor([0.1391], device='cuda:0', requires_grad=True)
```

DYNAMIC COMPUTATION GRAPH

Computation Graph (Chain Rule)

- Computational graphs are a type of graph that can be used to represent mathematical expressions
- These can be used for two different types of calculations:
 - Forward computation
 - Backward computation
- Static Vs Dynamic Computation Graph

Example




$$\frac{\partial Y}{\partial a} = \frac{\partial Y}{\partial d} \times \frac{\partial d}{\partial a} = e \times 1 = e$$

$$\frac{\partial Y}{\partial b} = \frac{\partial Y}{\partial d} \times \frac{\partial d}{\partial b} = e \times 1 = e$$

$$\frac{\partial Y}{\partial c} = \frac{\partial Y}{\partial e} \times \frac{\partial e}{\partial c} = d \times -1 = -d$$

Computation Graph (Chain Rule)




$$\hat{y} = z = x_1 w_1 + x_2 w_2$$

$$J = f(z, y)$$

$$dw = \frac{\partial J}{\partial w} = \frac{\partial J}{\partial z} * \frac{\partial z}{\partial w}$$

Computation Graph (Chain Rule)



PyTorch Computation Graph

- You have to see it for yourself
- The **PyTorchViz** package and its `make_dot(variable)`

```
from torchviz import make_dot
```

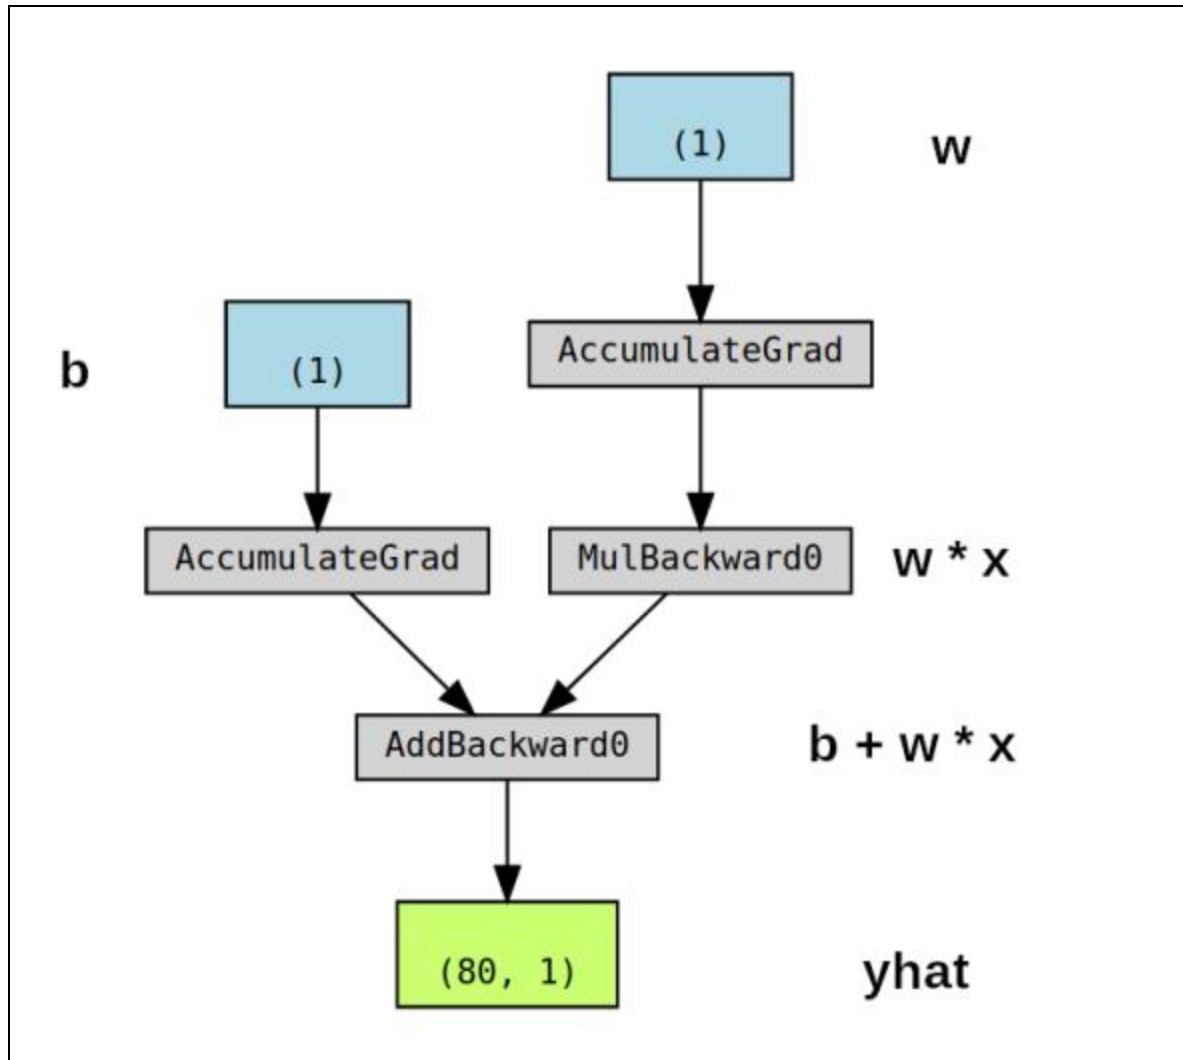
```
# Step 0 - Initializes parameters "b" and "w" randomly
torch.manual_seed(42)
b = torch.randn(1, requires_grad=True, \
                dtype=torch.float, device=device)
w = torch.randn(1, requires_grad=True, \
                dtype=torch.float, device=device)

# Step 1 - Computes our model's predicted output - forward pass
yhat = b + w * x_train_tensor

# Step 2 - Computes the loss
error = (yhat - y_train_tensor)
loss = (error ** 2).mean()

# We can try plotting the graph for any python variable:
# yhat, error, loss...
make_dot(yhat)
```

PyTorch Computation Graph



Where is X?

AUTOGRAD

Autograd

- Autograd is PyTorch's automatic differentiation package.
- Thanks to it, we don't need to worry about partial derivatives, chain rule, or anything like it.
- Role of the **backward()** method
 - It will compute gradients for all (requiring gradient) tensors involved in the computation of a given variable
- Do you remember the **starting point** for computing the gradients?
 - *loss.backward()*

Autograd in Action

```
yhat = b + w * x_train_tensor
```

```
error = (yhat - y_train_tensor)
```

```
loss = (error ** 2).mean()
```

```
# b_grad = 2 * error.mean()
```

```
# w_grad = 2 * (x_tensor * error).mean()
```

```
loss.backward()
```

Which tensors are going to be handled by the backward() method applied to the loss?

Autograd in Action

```
print(error.requires_grad, yhat.requires_grad, \
      b.requires_grad, w.requires_grad)
print(y_train_tensor.requires_grad, x_train_tensor.requires_grad)
```

```
True True True True
False False
```

grad

- What about the actual values of the gradients? We can inspect them by looking at the grad attribute of a tensor.

```
print(b.grad, w.grad)
```

```
tensor([-3.3881], device='cuda:0')  
tensor([-1.9439], device='cuda:0')
```

Gradient Accumulation

- If you check the method's documentation, it clearly states that **gradients are accumulated**
- **Why?**
- If we run **above** code twice and check the grad attribute afterward

```
tensor([-6.7762], device='cuda:0')  
tensor([-3.8878], device='cuda:0')
```

Zeroing Accumulation

- Every time we use the gradients to update the parameters, we need to zero the gradients afterward. And that's what **zero_()** is good for

```
b.grad.zero_(), w.grad.zero_()
```

```
(tensor([0.], device='cuda:0'),  
 tensor([0.], device='cuda:0'))
```

no_grad

- Perform regular Python operations on tensors, without affecting PyTorch's computation graph
- Will be used during inference

```
with torch.no_grad():  
    b -= lr * b.grad  
    w -= lr * w.grad
```


Summary So Far

```
b = torch.randn(1, requires_grad=True, \
                dtype=torch.float, device=device)
```

```
loss = (error ** 2).mean()
```

```
loss.backward()
```

```
print(b.grad, w.grad)
```

```
b.grad.zero_(), w.grad.zero_()
```


```
with torch.no_grad():
    b -= lr * b.grad
    w -= lr * w.grad
```

OPTIMIZERS

Optimizer

- We've been manually updating the parameters using the computed gradients so far
- What if we had a whole lot of parameters?!
- We need to use one of PyTorch's optimizers, like **SGD**, RMSprop, or Adam.
- SGD is the most basic of them, and Adam is one of the most popular.
- Different optimizers use different mechanics for updating the parameters, but they all achieve the same goal through, literally, different paths

Optimizer



- Remember, the choice of mini-batch size influences the path of gradient descent, and so does the choice of an optimizer

Optimizer (step / zero_grad)

- An optimizer takes the parameters we want to update, the **learning rate** we want to use (and possibly many other **hyper-parameters** as well!), and performs the updates through its **step()** method

```
import torch.optim as optim
```

```
# Defines an SGD optimizer to update the parameters  
optimizer = optim.SGD([b, w], lr=lr)
```

```
#      b -= lr * b.grad  
#      w -= lr * w.grad  
optimizer.step()
```

```
# b.grad.zero_  
# w.grad.zero_  
optimizer.zero_grad()
```

LOSS FUNCTIONS IN PYTORCH

Loss functions in PyTorch

- PyTorch got us covered with many loss functions to choose from, depending on the task at hand.
- Since ours is a regression, we are using the Mean Squared Error (MSE) as loss, and thus we need PyTorch's `nn.MSELoss`

```
# Defines a MSE loss function  
loss_fn = nn.MSELoss(reduction='mean')
```

```
predictions = torch.tensor(0.5, 1.0)  
labels = torch.tensor(2.0, 1.3)  
loss_fn(predictions, labels)
```

```
tensor(1.1700)
```

Loss functions in PyTorch

```
# Defines a MSE loss function  
loss_fn = nn.MSELoss(reduction='mean')
```

```
loss.cpu().numpy()
```

```
# error = (yhat - y_train_tensor)  
# loss = (error ** 2).mean()  
loss = loss_fn(yhat, y_train_tensor)
```

```
loss.detach().cpu().numpy()
```

```
# Step 3 - Computes gradients for both "b" and "w" parameters  
loss.backward()
```

```
# Step 4 - Updates parameters using gradients and  
# the learning rate  
optimizer.step()  
optimizer.zero_grad()
```

```
print(loss.item(), loss.tolist())
```


PYTORCH MODEL

Model Class

- Are you comfortable with object-oriented programming (OOP) concepts like classes, constructors, methods, instances, and attributes?
 - <https://realpython.com/python3-object-oriented-programming/>
 - Having a good understanding of OOP is key to benefit the most from PyTorch's capabilities.

Model Class Parts

- Must inherit from **nn.Module**
- **__init__(self):**
 - it defines the parts that make up the model
 - In our case, two parameters, b and w.
 - You are **not** limited to defining **parameters**, though... **models can contain other models as their attributes** as well, so you can easily nest them. We'll see an example of this shortly as well
 - Do not forget to include **super().__init__()** to execute the **__init__()** method of the parent class (nn.Module) before your own

Model Class Parts

- `forward(self, x):`
 - It performs the actual computation, that is, it outputs a **prediction, given the input x**
 - We use `model(x)` instead of `forward(x)`

Model Class Parts

- `forward(self, x):`
 - It performs the actual computation, that is, it outputs a **prediction, given the input x**
 - We use `model(x)` instead of `forward(x)`

Our Regression Model

```
class ManualLinearRegression(nn.Module):
    def __init__(self):
        super().__init__()
        # To make "b" and "w" real parameters of the model,
        # we need to wrap them with nn.Parameter
        self.b = nn.Parameter(torch.randn(1,
                                           requires_grad=True,
                                           dtype=torch.float))

        self.w = nn.Parameter(torch.randn(1,
                                           requires_grad=True,
                                           dtype=torch.float))

    def forward(self, x):
        # Computes the outputs / predictions
        return self.b + self.w * x
```

Model Parameters

- We define our two parameters, b and w , using the **Parameter() class**
- This tells PyTorch that these **tensors**, which are **attributes** of the ManualLinearRegression **class**, should be considered **parameters of the model**
- we can use our model's **parameters()** method to retrieve an iterator over all model's parameters, including parameters of nested model
- we can use it to feed our optimizer (instead of building a list of parameters ourselves!)

Model `state_dict`

- We can get the current values of all parameters using our model's `state_dict()` method
- Only learnable parameters are included, as its purpose is to keep track of parameters

```
dummy.state_dict()
```

```
OrderedDict([('b', tensor([0.3367])), ('w', tensor([0.1288]))])
```


Model Device

- We need to send our model to the same device where data is
 - If our data is made of GPU tensors, our model must “live” inside the GPU as well

```
torch.manual_seed(42)
# Creates a "dummy" instance of our ManualLinearRegression model
# and sends it to the device
dummy = ManualLinearRegression().to(device)
```

Nested Models

- PyTorch's Linear model/layer

```
linear = nn.Linear(1, 1)  
linear
```

```
Linear(in_features=1, out_features=1, bias=True)
```

```
linear.state_dict()
```

```
OrderedDict([('weight', tensor([[[-0.2191]]])),  
            ('bias', tensor([0.2018]))])
```

Nested Models

```
class MyLinearRegression(nn.Module):
    def __init__(self):
        super().__init__()
        # Instead of our custom parameters, we use a Linear model
        # with a single input and a single output
        self.linear = nn.Linear(1, 1)

    def forward(self, x):
        # Now it only takes a call
        self.linear(x)
```

```
dummy = MyLinearRegression().to(device)
list(dummy.parameters())
```

```
dummy.state_dict()
```

```
OrderedDict([('linear.weight',
                  tensor([[0.7645]], device='cuda:0')),
              ('linear.bias',
                  tensor([0.8300], device='cuda:0'))])
```

Summary So Far

- Optimizer
 - `opt=optim.SGD(parameters)`
 - `opt.step()`
 - `opt.zero_grad()`
- Loss Fuction
 - `nn.MSELoss()`
 - `loss.backward()`
- Model Class
 - `__init__(self)`
 - `forward(self, x)`
 - Defining Model Parameters
 - `Model state_dict()`
 - Model device
 - Nested Model

A SIMPLE REGRESSION PROBLEM (PYTORCH IMPLEMENTATION)

Imports, Device, Seed

```
1  import numpy as np
2  import matplotlib.pyplot as plt
3  import torch
4  from torchviz import make_dot
```

```
7  device='cuda' if torch.cuda.is_available() else 'cpu'
8  torch.manual_seed(100)
```

Model Class

```
10 class SimpleRGNet(torch.nn.Module):
11     def __init__(self):
12         super().__init__()
13         self.linear = torch.nn.Linear(1,1,bias=True)
14     def forward(self,x):
15         return self.linear(x)
16
```

Model, optimizer and loss initialization

```
40 model=SimpleRGNet().to(device)
41 paramList=list(model.parameters())
42 stateDict=model.state_dict()
43 print(paramList)
44 print(stateDict)
45
46 lr=0.1
47 optimizer=torch.optim.SGD(model.parameters(),lr=lr)
48 lossfnc = torch.nn.MSELoss(reduce="mean")
```


Data Preparation

```
20 true_w=2
21 true_b=1
22 N=100
23 #data generation
24 np.random.seed(100)
25 x=np.random.rand(N,1)
26 epsilon=0.1*np.random.randn(N,1)
27 y=true_w*x+true_b+epsilon
28 #data split
29 idx=np.arange(N)
30 np.random.shuffle(idx)
31 idx_train=idx[:int(0.8*N)]
32 idx_test=idx[int(0.8*N):]
33 x_train, y_train = x[idx_train],y[idx_train]
34 x_val, y_val = x[idx_test],y[idx_test]
35 x_train_tensor=torch.as_tensor(x_train).float().to(device)
36 y_train_tensor=torch.as_tensor(y_train).float().to(device)
37 x_val_tensor=torch.as_tensor(x_val).float().to(device)
38 y_val_tensor=torch.as_tensor(y_val).float().to(device)
```

Training Loop

```
53 trainLosses=[]
54 valLosses=[]
55 for i in range(1000):
56     model.train()
57     #forward pass
58     yhat=model(x_train_tensor)
59     loss=lossfnc(yhat,y_train_tensor)
60     trainLosses.append(loss.item())
61     #make_dot(loss).view()
62     loss.backward()
63     optimizer.step()
64     optimizer.zero_grad()
65     stateDict=model.state_dict()
66     w=stateDict['linear.weight']
67     b=stateDict['linear.bias']
68     w=w.item()
69     b=b.item()
70
71     model.eval()
72     with torch.no_grad():
73         #val MSE loss
74         yhatval=model(x_val_tensor)
75         valLoss=lossfnc(yhatval,y_val_tensor)
76         valLosses.append(valLoss.item())
77     #stopping condition
78     if(valLoss.item()<0.0001):
79         break
80     print(f'train loss={loss.item()}, val loss={valLoss.item()}, w={w}, b={b}')
```

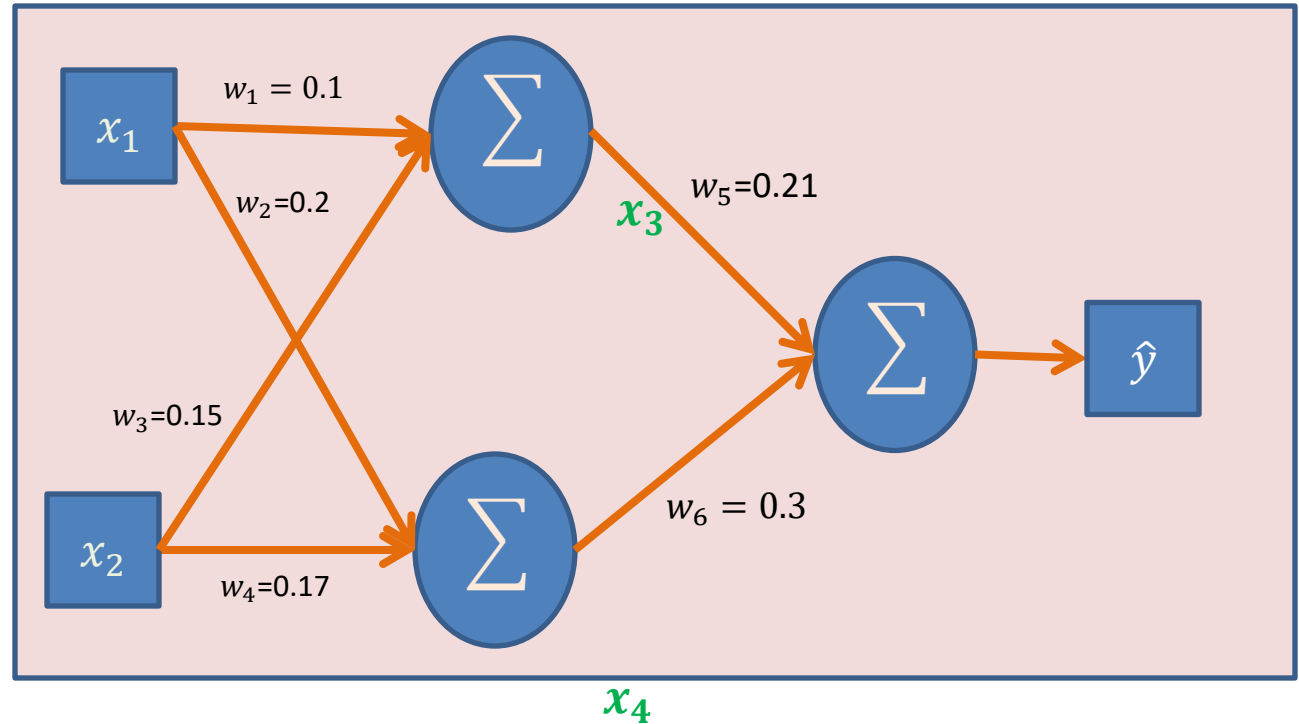
PyTorch Summary

- Data & Parameters as Tensors
- Moving from numpy to Tensor to GPU
- PyTorch Model Class
- Autograd using backward()
- Built-in Loss Functions
- Built-in optimizers
- Much Simplified Training Loop
- Getting Ready for deep models training using large datasets?

Graded Home Task (backprop)

Training Data

x_1	x_2	y
-2	3	5
1	-2	-3



- Perform weight update for one epoch using batch gradients and $\text{lr}=0.1$
- Do all steps manually first
- Write a PyTorch code using Autograd to verify your results
- Submit both manual calculation sheet & PyTorch code