

Google Classroom Code: mhxgl24

Basics of Neural Networks

Deep Learning (DS-5006)

Dr. Adeel Mumtaz

Lecture 3

Fall, 2022




National University
Of Computer and Emerging Sciences

Contents

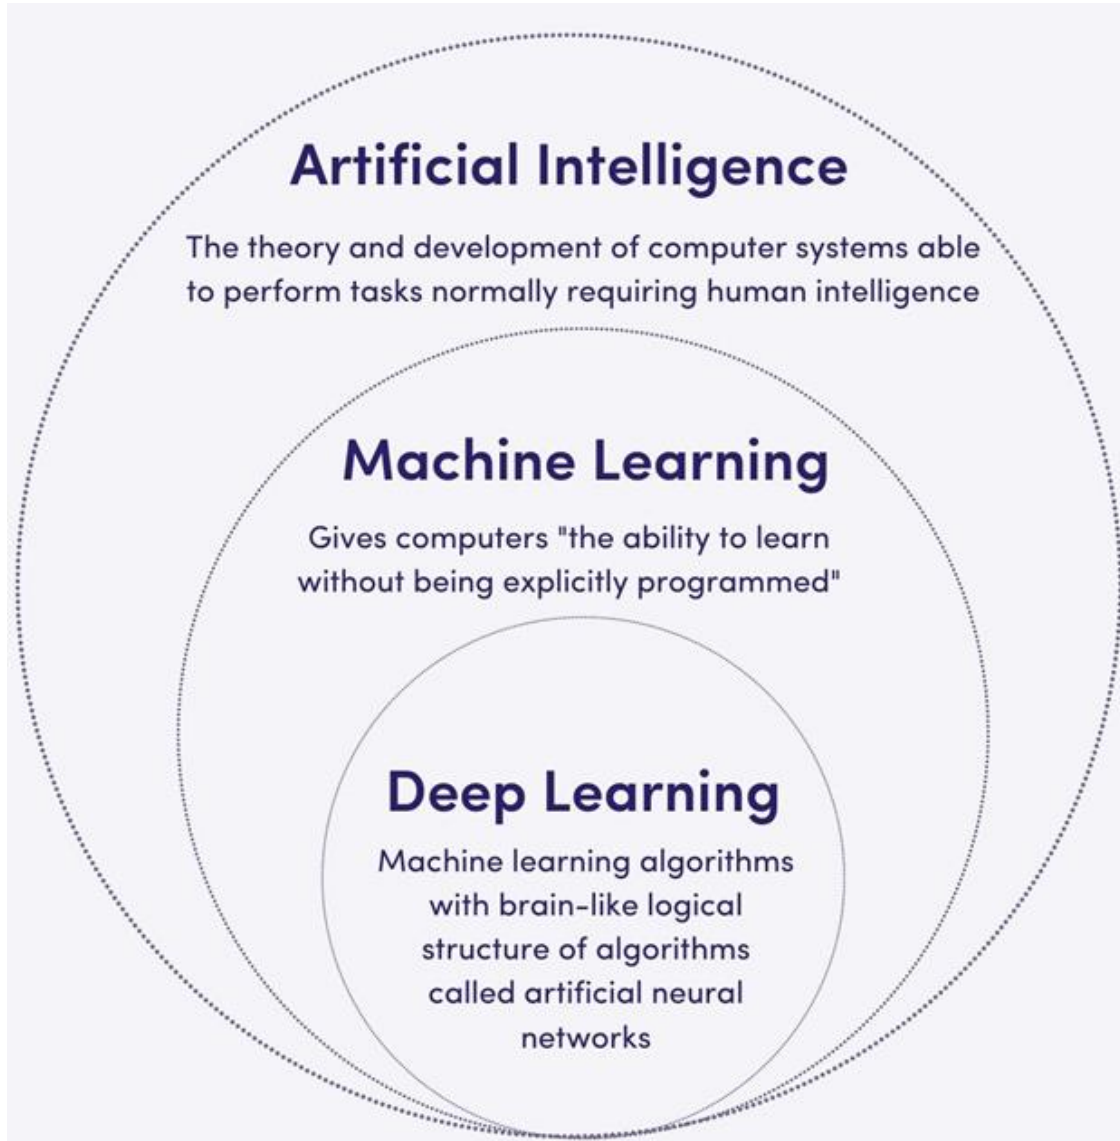
- Basics of Neural Networks
 - Deep Learning
 - BNN Vs ANN
 - An Intuitive example (housing price prediction)
 - Structured and Unstructured data
 - Shallow Vs Deep Neural Networks
 - Deep Learning architectures
- A Simple Regression Problem (Theory)
 - Network Architecture
 - MSE Loss Function
 - Gradient Descent Algorithm
 - Learning Rate Effect
 - Training/Inference Loop
 - **Batch vs Stochastic vs Mini-Batch GD**
- A Simple Regression Problem (Numpy Implementation)
 - Data Generation
 - Data Splitting
 - Visualizing Data
 - Parameter initialization
 - Training Loop
 - Loss Calculation
 - Gradient Calculations
 - Parameter Updates
 - Validation Loss / Stopping Criteria
 - Plots/Training Curves
- NN Summary
- Home Task

BASICS OF NEURAL NETWORKS

Deep Learning




Deep Learning



- The term Deep Learning refers to **training very large Neural Network**


What is a Neural Network

- **Biologically** motivated approach to machine learning
- **Artificial** neural network (**ANN**) is a machine learning approach that models **human brain** and consists of a number of artificial **neurons**
- Can be used for
 - Classification
 - Regression
 - Clustering
 - Association
 - Optimization



Biological Neural Networks

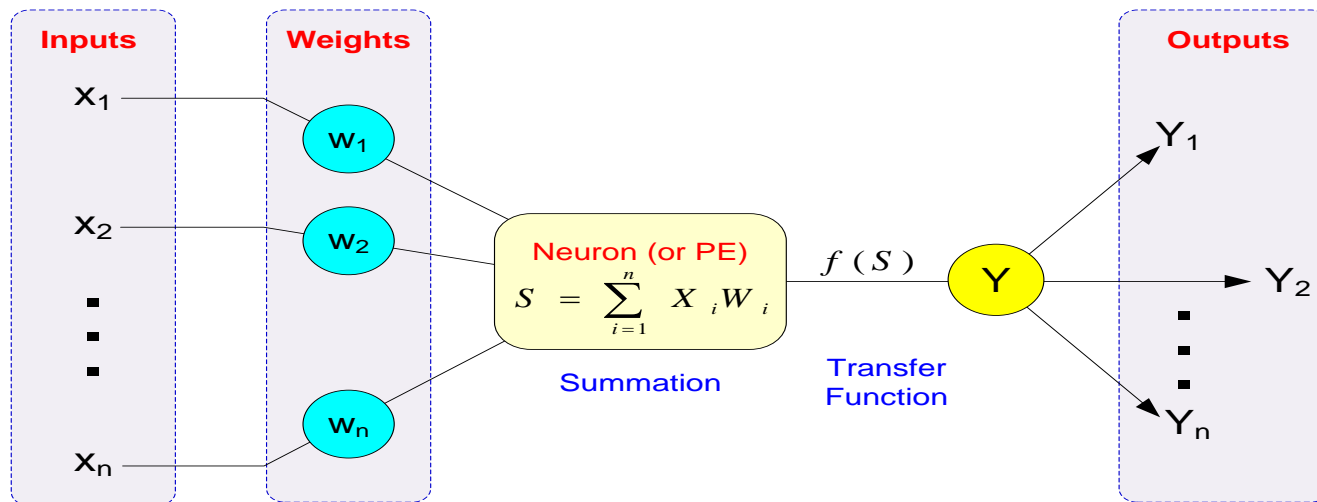
- Fundamental processing element of a neural network is a **neuron**
 - Dendrites, accept inputs
 - Soma, process the inputs
 - Axon, turns the processed inputs into outputs
 - Synapses, the electrochemical contacts between neurons
- A human brain has 100 billion neurons
- An ant brain has 250,000 neurons



Two interconnected brain cells (neurons)

Artificial Neural Networks

- A single neuron (processing element – PE) with inputs and outputs
 - Inputs
 - Weights
 - Summation S or V
 - Transfer function or Activation function, φ or f



Comparison between ANN & BNN


Biological	versus	Artificial	NNs
Soma		Node	
Dendrites		Input	
Axon		Output	
Synapse		Weight	
Slow		Fast	
Many neurons (10^9)		Few neurons (<u>~ 100s</u>)	

An Intuitive example (housing price prediction)

- Fit a function to predict the price of the house as a function of the size
 - linear regression?

House Size (X)	Price (Y)
500	5 Laks
600	8 Laks
700	9 Laks
800	13 Laks
900	12 Laks
1100	18 Laks
1200	20 Laks


Housing Price Prediction



An Intuitive example (housing price prediction)

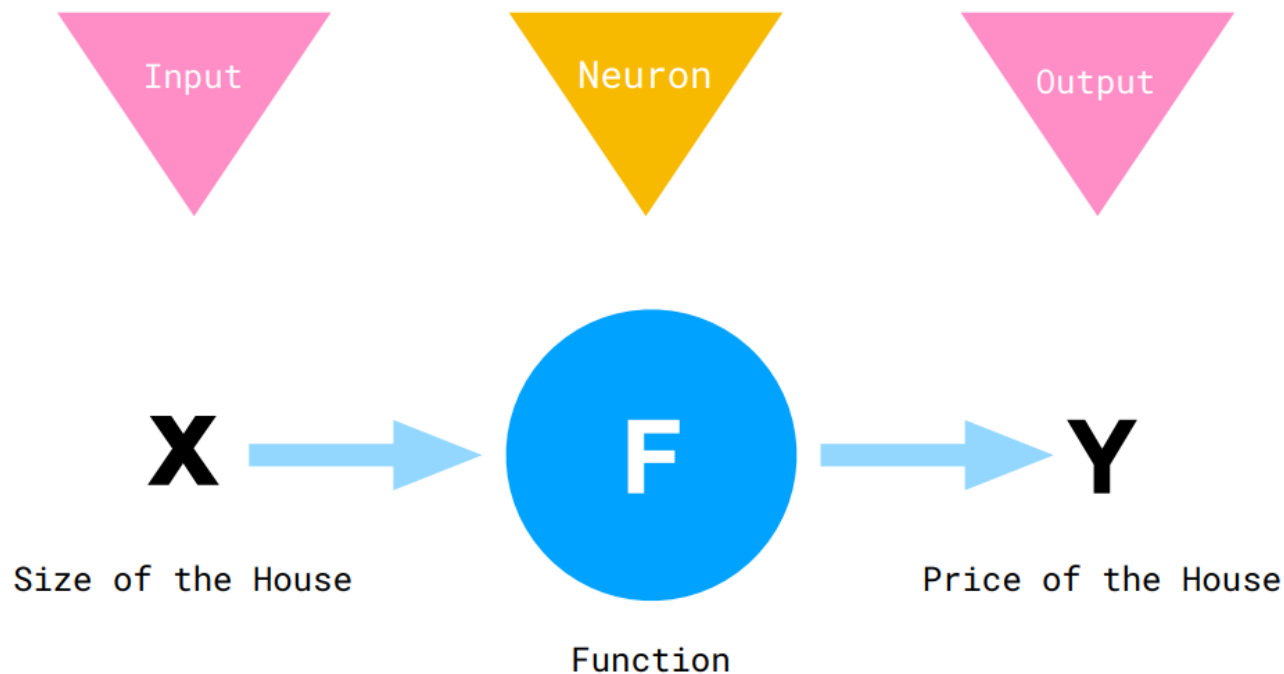
- Price can't be negative

ReLU Function



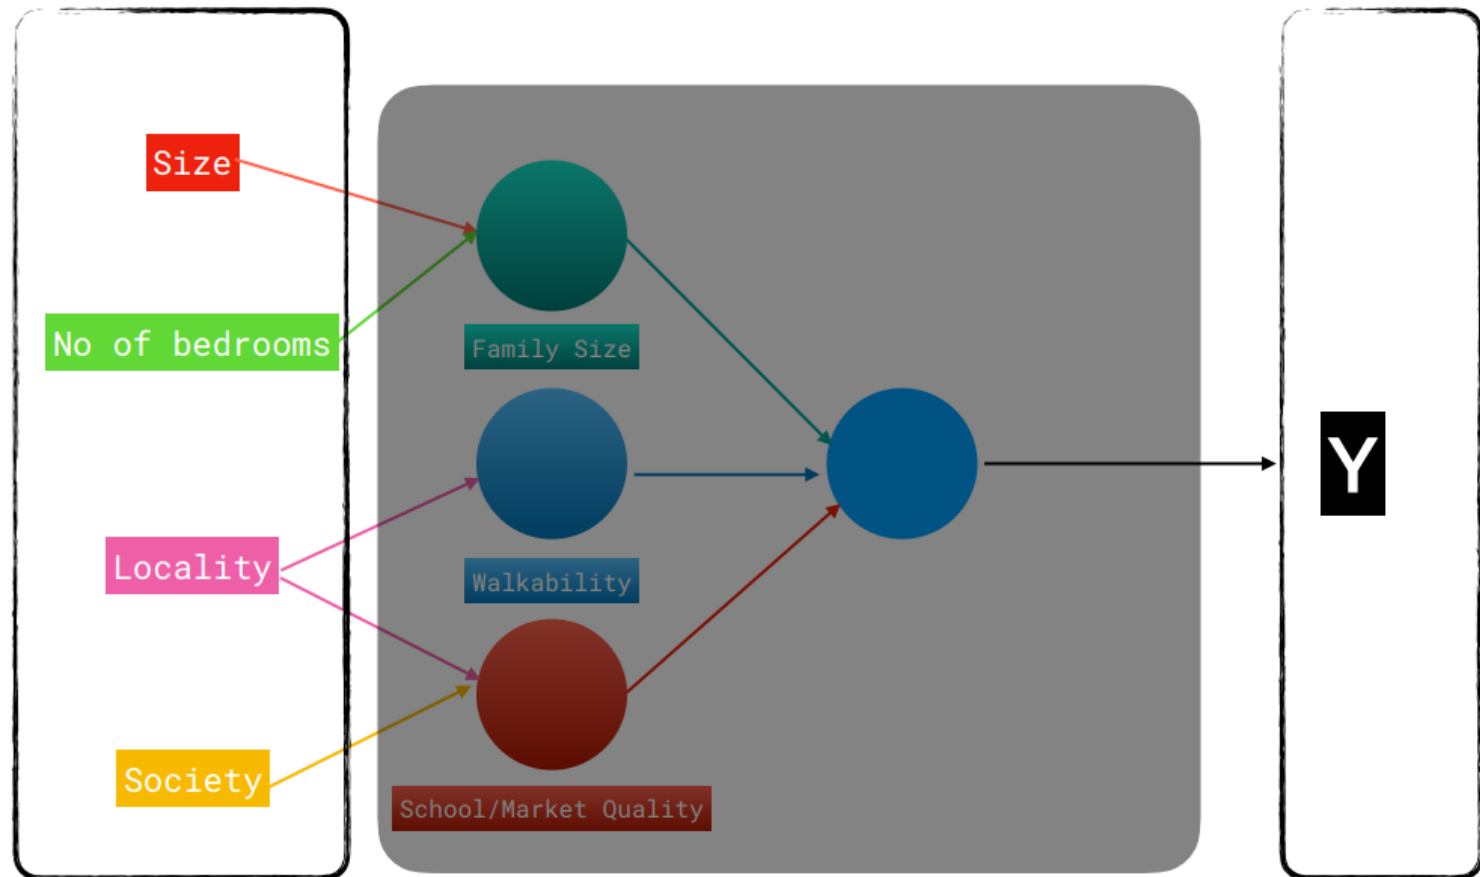
An Intuitive example (housing price prediction)

- a very simple neural network
- Neuron takes inputs the size, computes the linear function, and then outputs the estimated price.




An Intuitive example (housing price prediction)

- Other features that can affect the Price



An Intuitive example (housing price prediction)

- A neural network with four inputs and one output
- One hidden layer with three hidden units
- Note that each of the hidden unit takes its inputs of all four features (Dense, Linear, Fully connected)



Structured Data

Size	#bedrooms	...	Price (1000\$s)
1200	2		300
1500	3		400
2000	3		480
⋮	⋮		⋮
3000	4		520

User Region	Ad Id	...	Ad revenue(\$)
USA	1005		0.5
UK	1009		0.3
USA	998		0.5
⋮	⋮		⋮
CAN	2104		0.45

Unstructured Data



Audio











Image

Once upon a time in
the history of...

Text


Unstructured Data

 Text files and documents	 Server, website and application logs	 Sensor data	 Images
 Video files	 Audio files	 Emails	 Social media data


Deep Neural Networks are very good at processing these Unstructured data

Shallow Vs Deep Neural Networks

SHALLOW NEURAL NETWORK



DEEP NEURAL NETWORK




Shallow Vs Deep Neural Networks

Factors	Shallow Neural Network (SNN)	Deep Neural Network (DNN)
Feature Engineering	1. Individual feature extraction process is required. Various features cited in the literature are histogram oriented gradients, speeded up robust features, and local binary patterns.	1. Replace the hand-crafted features and directly work on the entire input. Thus, more practical for complex datasets.
Data Size Dependency	2. Needs a lesser quantity of data.	2. Needs vast volumes of data.

Deep Learning Architectures

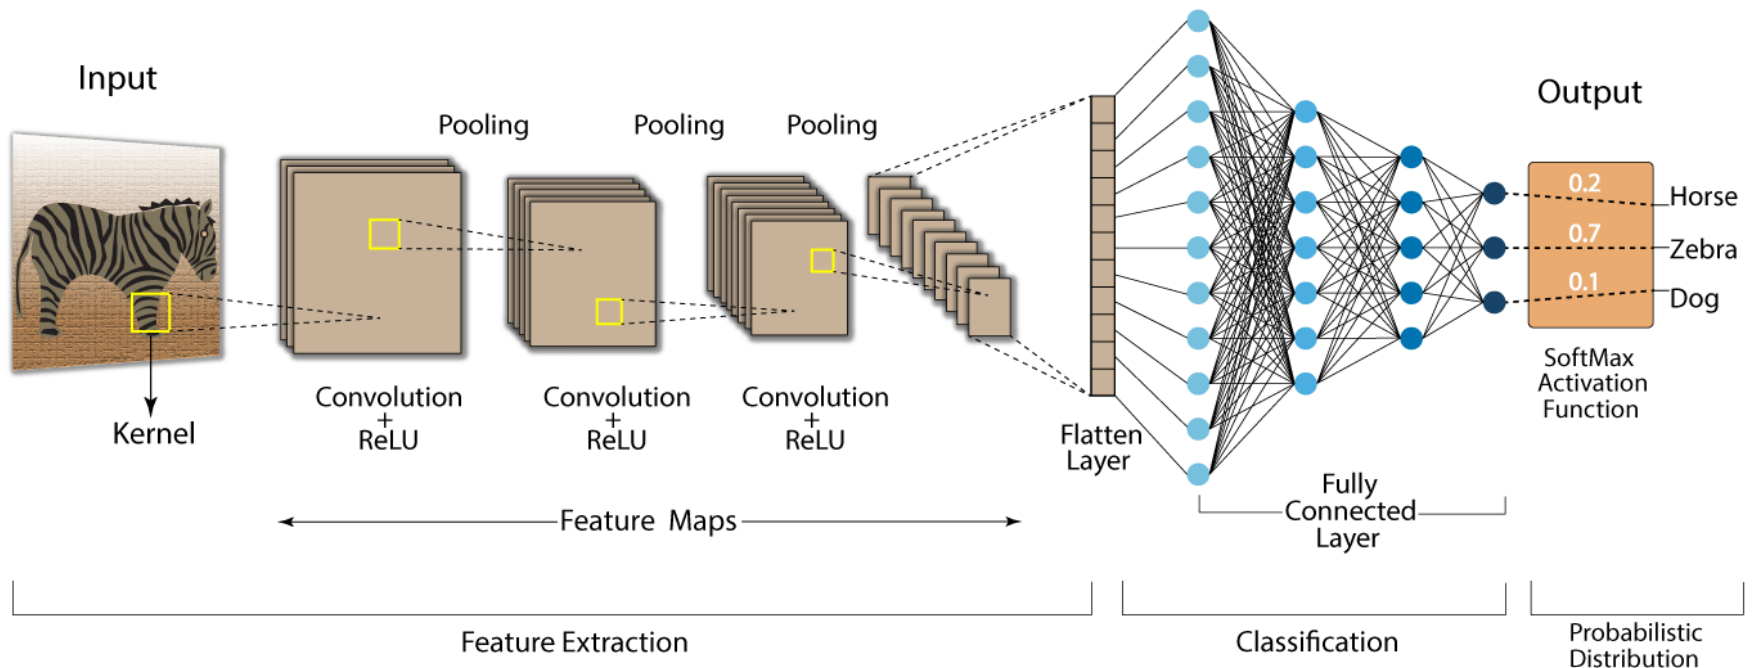
- Universal Function Approximators
- The number of trainable parameters increases drastically with an increase in the size of the image



Artificial Neural Network (ANN or MLP)

Deep Learning Architectures


- CNN learns the filters automatically
- CNN captures the spatial features from an image, identify object location accurately
- CNN follows the concept of parameter sharing



Convolutional Neural Network (CNN)

Deep Learning Architectures

- RNN captures the sequential information present in the input data
- RNNs share the parameters across different time steps

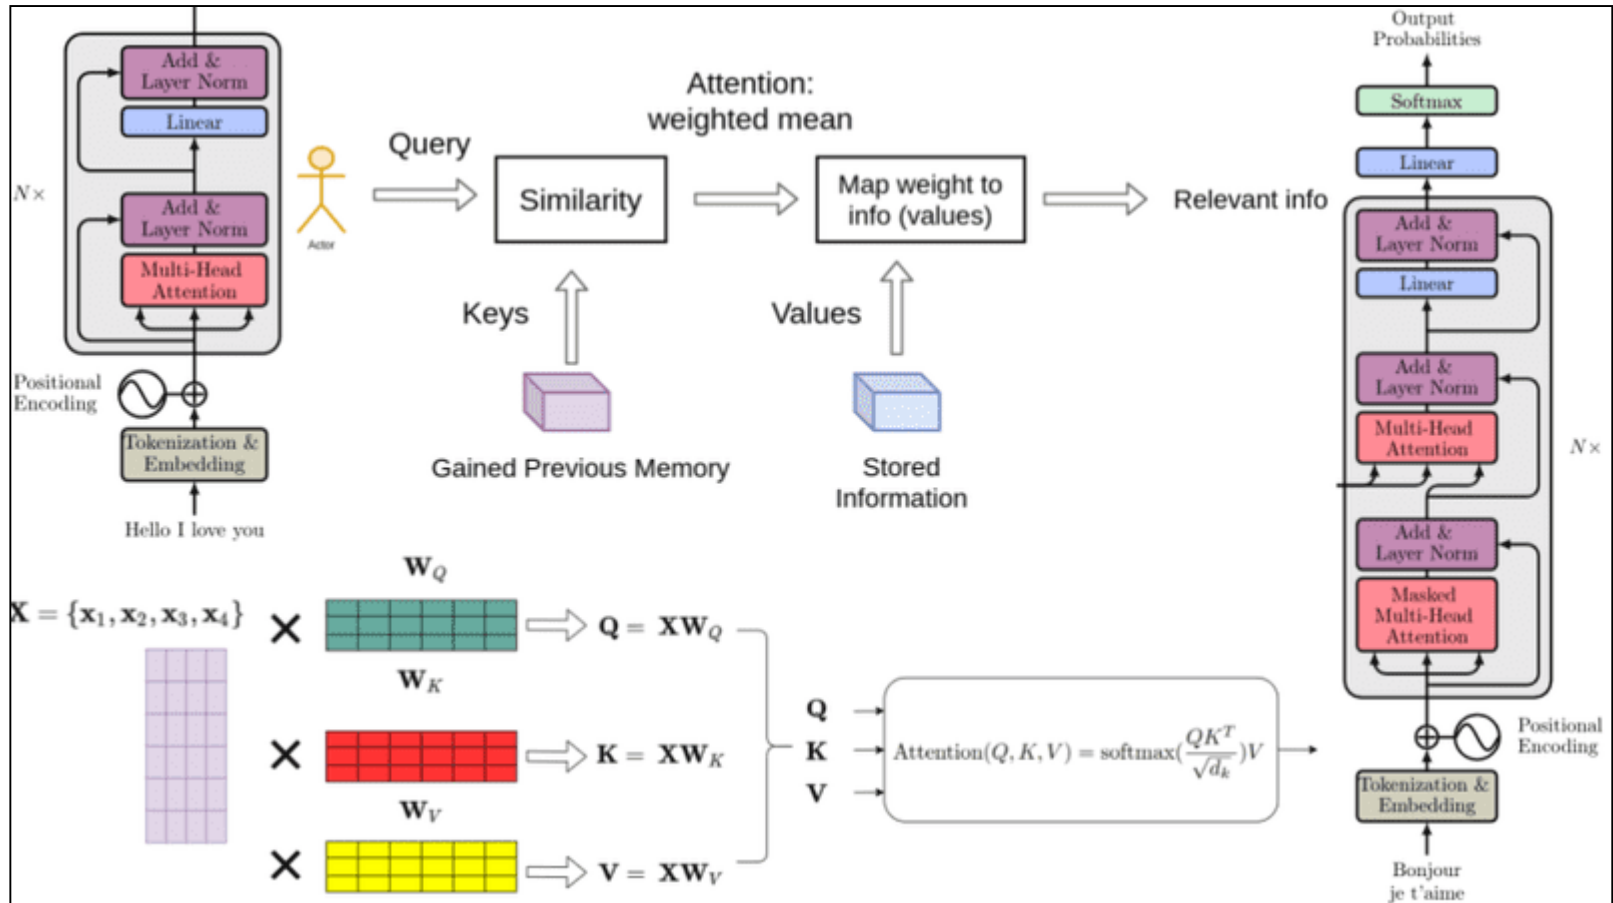


Recurrent Neural Network (RNN)

Deep Learning Architectures

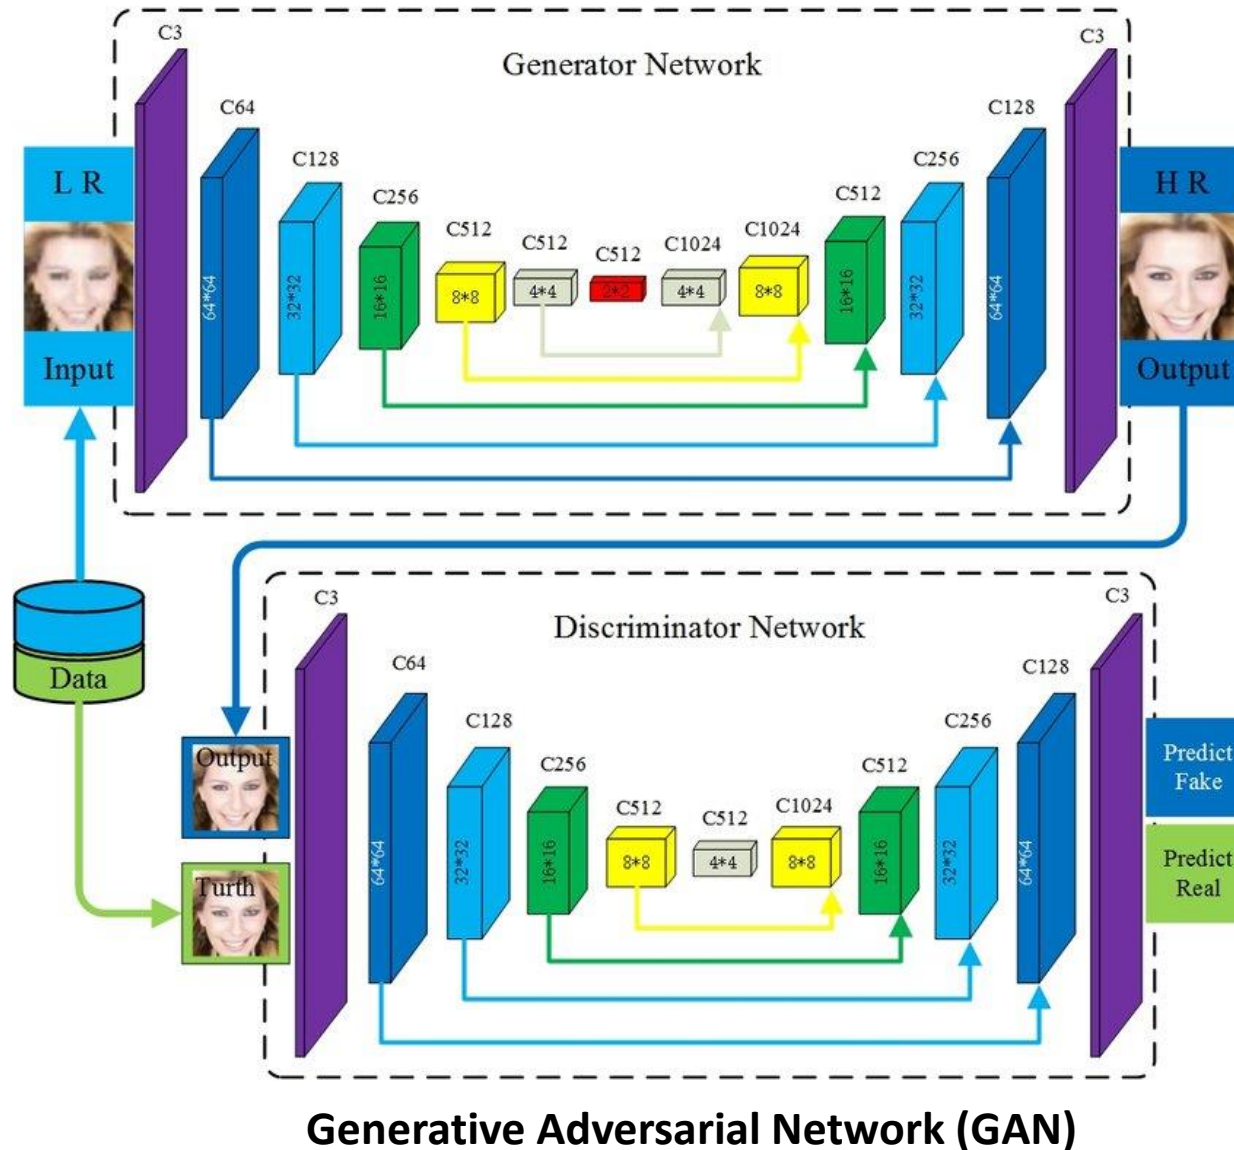
	MLP	RNN	CNN
Data	Tabular data	Sequence data (Time Series, Text, Audio)	Image data
Recurrent connections	No	Yes	No
Parameter sharing	No	Yes	Yes
Spatial relationship	No	No	Yes
Vanishing & Exploding Gradient	Yes	Yes	Yes

Deep Learning Architectures



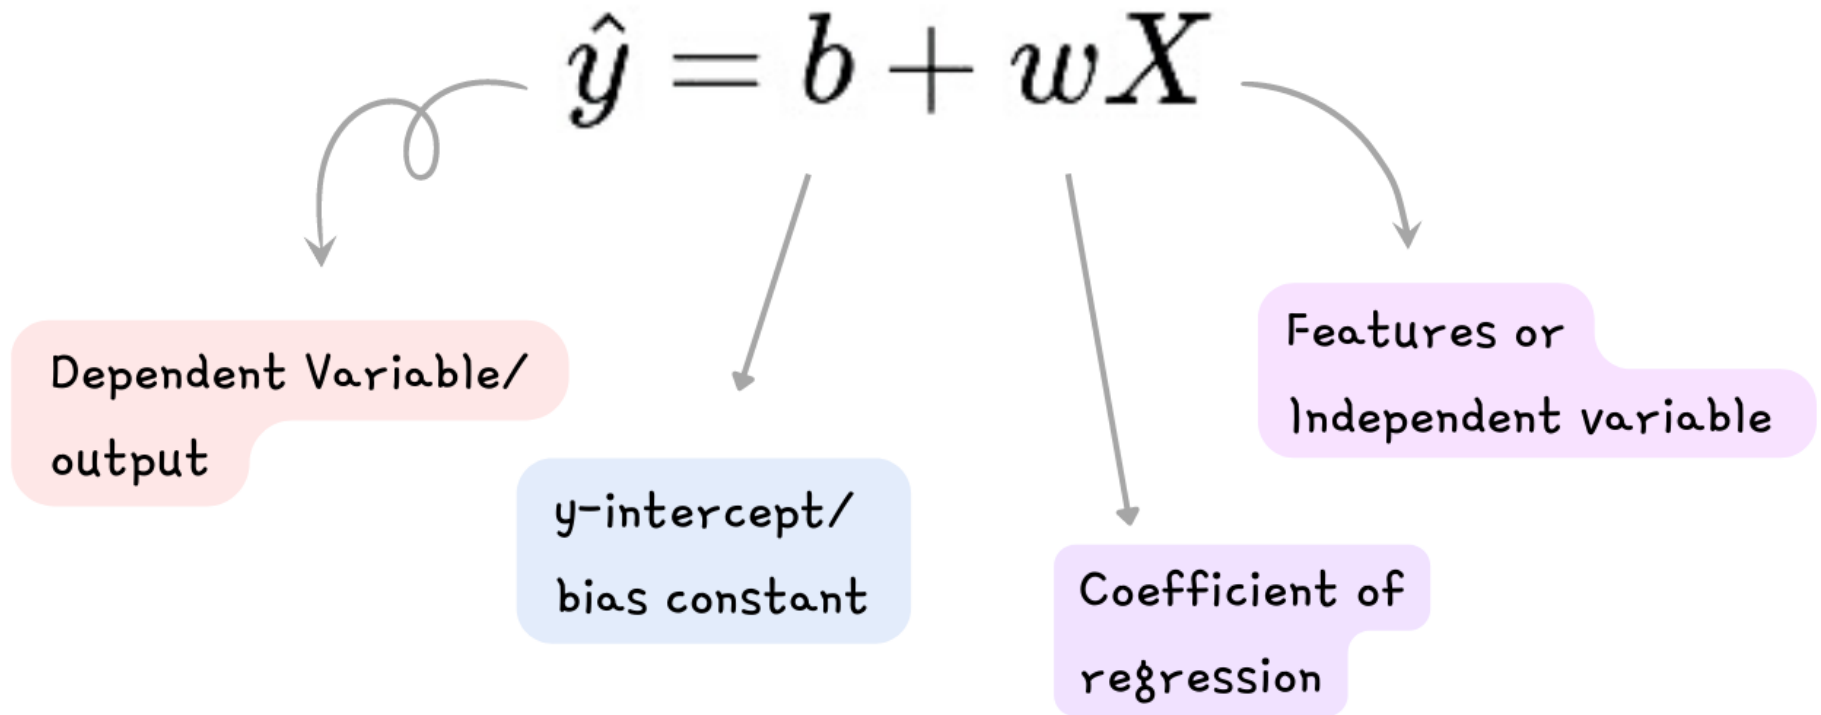
Attention Based Neural Network (Transformer)

Deep Learning Architectures




A SIMPLE REGRESSION PROBLEM (THEORY)

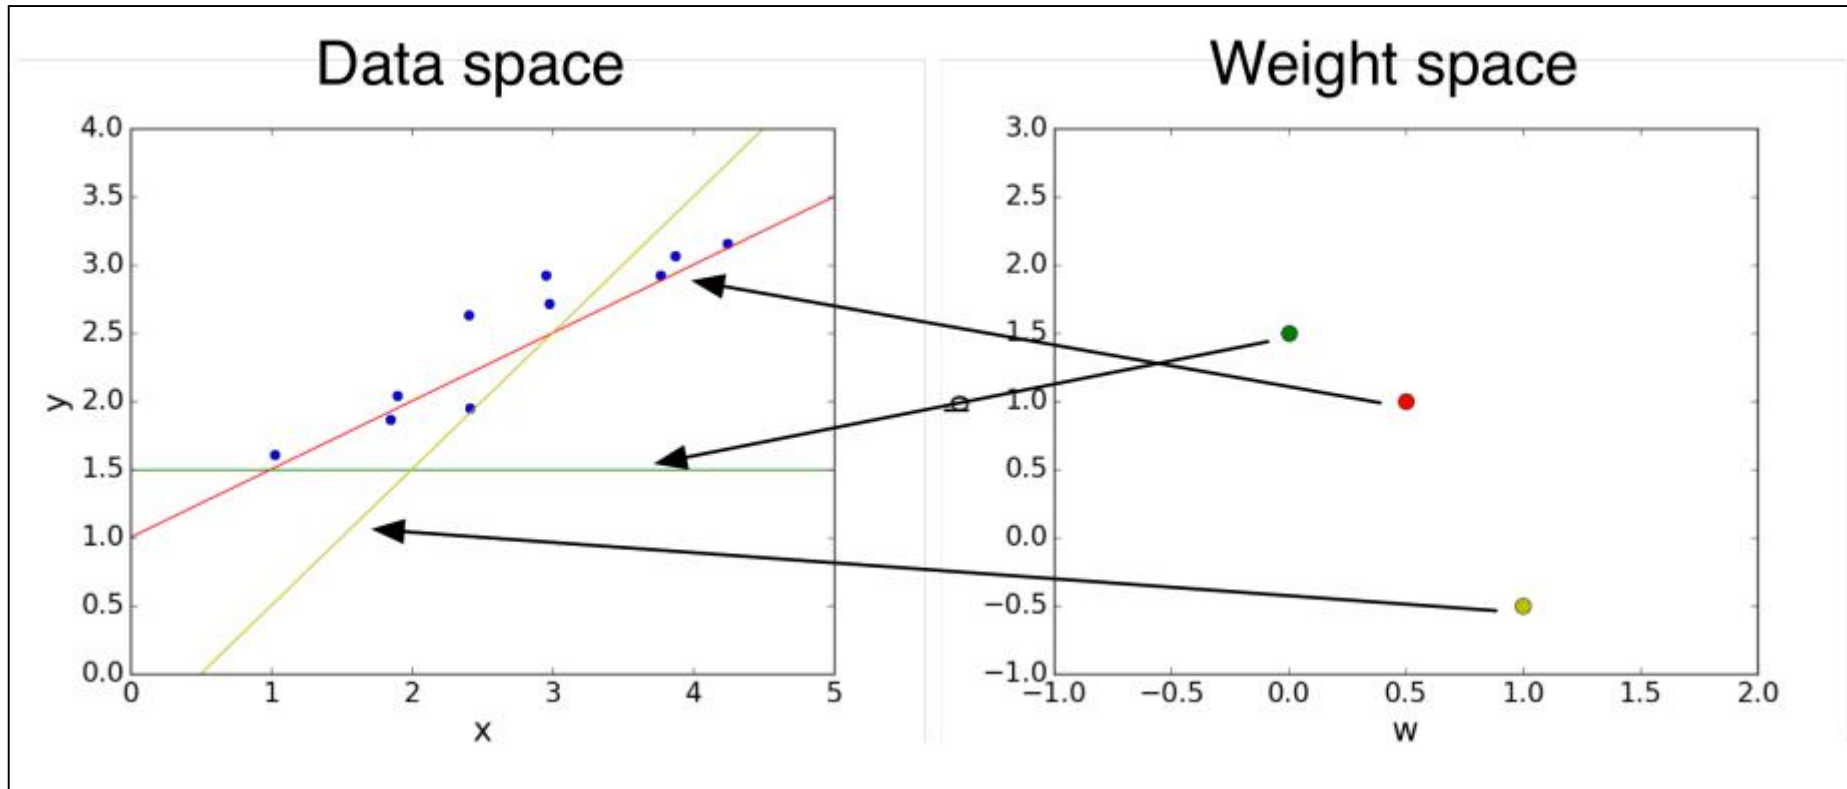
Simple Linear Regression



Simple Linear Regression



Simple Linear Regression




Simple Linear Regression


$$Cost(b, w) = \mathcal{L}(b, w) = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

$$= \mathcal{L}(x, b, w) = \frac{1}{N} \sum_{i=1}^N (wx_i + b - y_i)^2$$

$Cost(b, w)$




Simple Linear Regression



The Linear Unit: $y = wx + b$


- **NN=Architecture + Parameters**
- **Training NN = Given data learn best Parameters which gives minimum loss (usually an iterative process/algorithm)**

Simple Linear Regression



Gradient Descent Algorithm

$$W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$$



Simple Linear Regression

$$\mathcal{L}(x, b, w) = \frac{1}{N} \sum_{i=1}^N (wx_i + b - y_i)^2$$

$$dw = \frac{\partial \mathcal{L}(x, b, w)}{\partial w} = 2 * \frac{1}{N} \sum_{i=1}^N (wx_i + b - y_i)(x_i)$$


$$= 2 * \frac{1}{N} \sum_{i=1}^N (error_i)(x_i) = 2 * mean(error * x)$$

$$db = \frac{\partial \mathcal{L}(x, b, w)}{\partial b} = 2 * \frac{1}{N} \sum_{i=1}^N (error_i) = 2 * mean(error)$$


$$w_{new} = w_{old} - \alpha * dw$$

$$b_{new} = b_{old} - \alpha * db$$

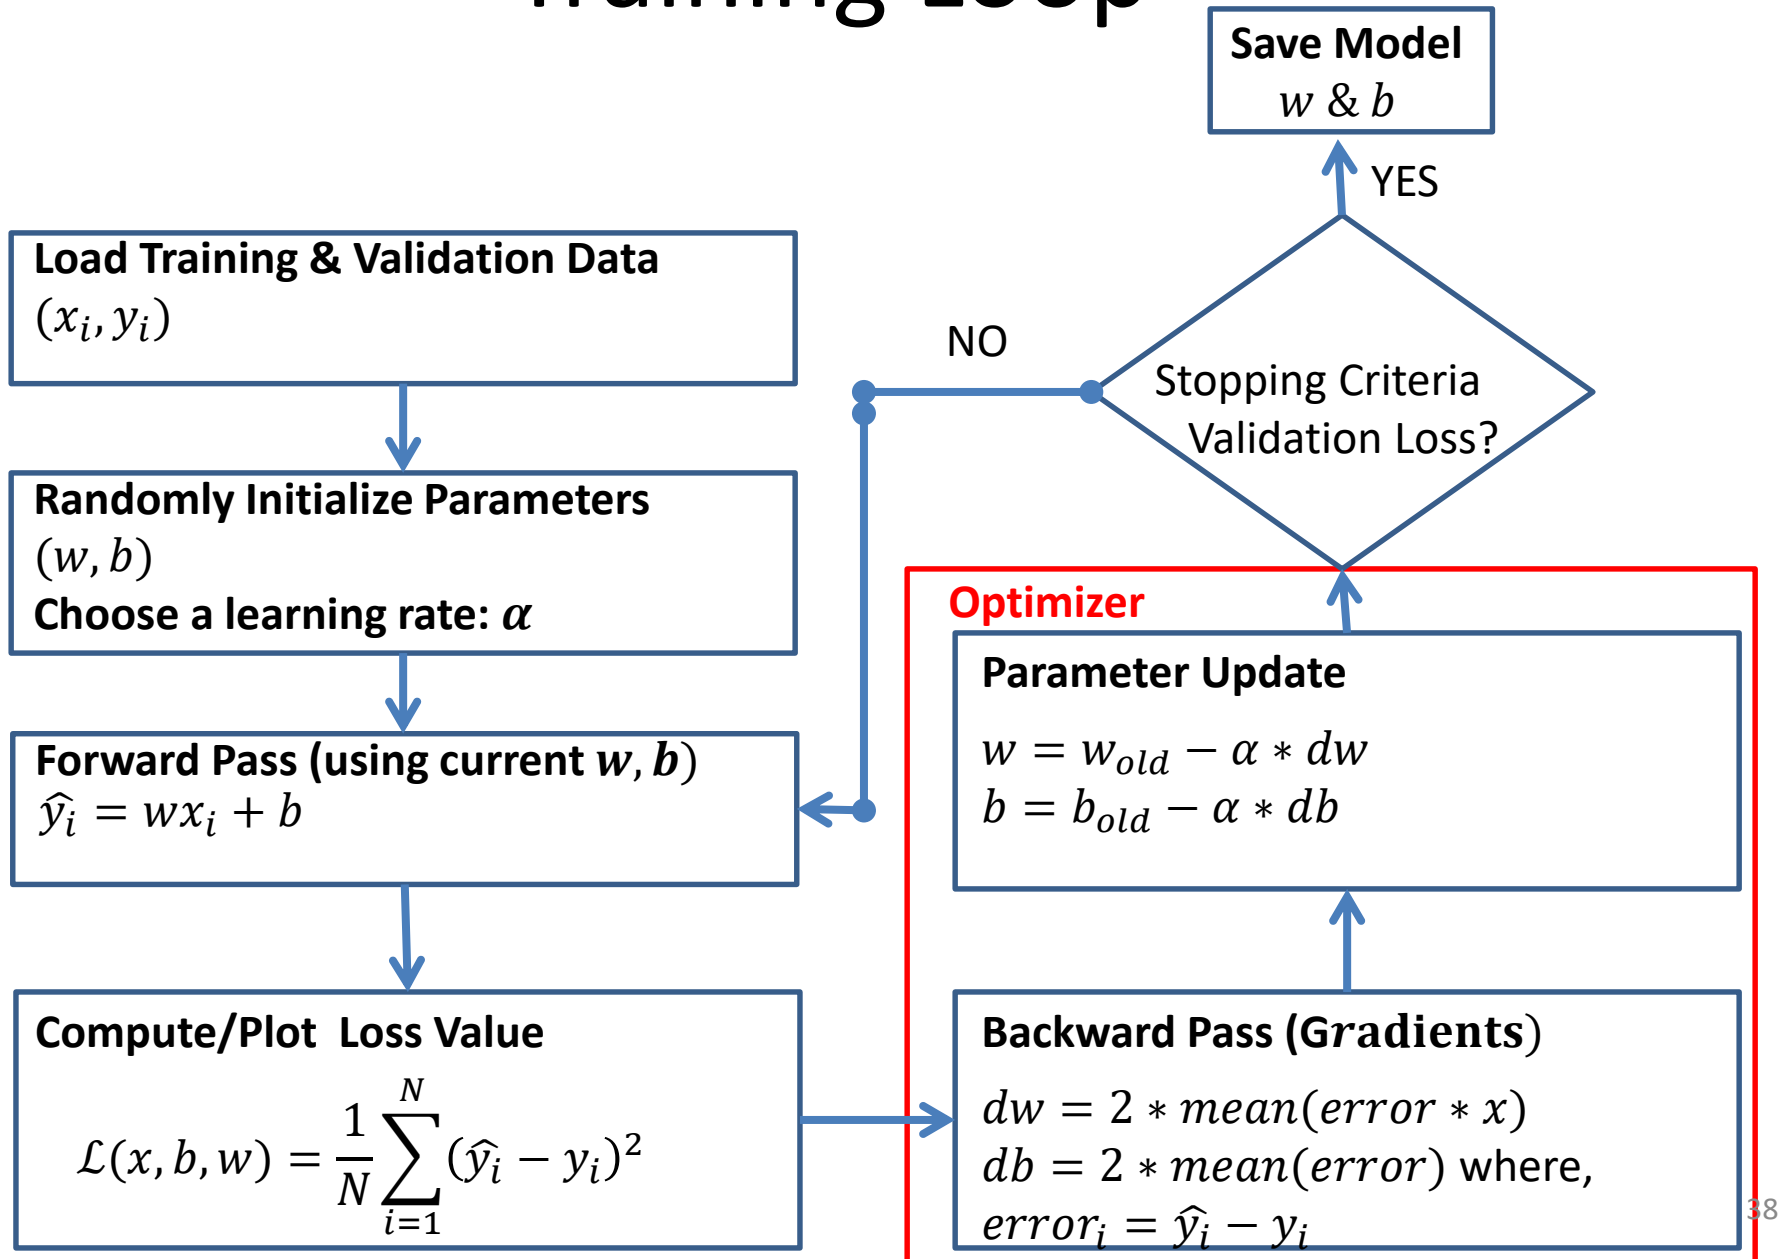
Effect of Learning Rate



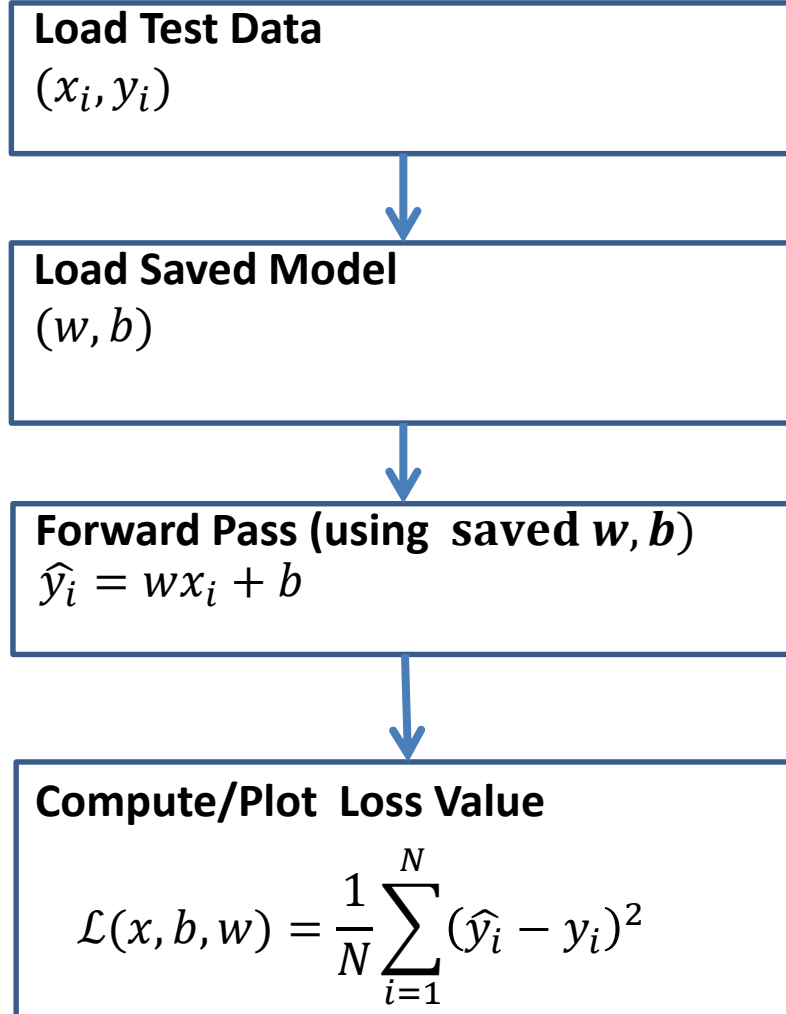
Effect of Learning Rate



Training Loop



Inference/Testing



Batch Gradient Descent

- An **epoch** refers to one cycle through the full training dataset
- All the training data is taken into consideration to update parameters
- Uses **mean of gradients** to update parameters
- One **step/iteration** of gradient descent in one epoch
- Great for **convex** or relatively smooth error surfaces
- The graph of **cost vs epochs** is also quite smooth


Stochastic Gradient Descent

- If our dataset has 5 million examples, then just to take one step the model will calculate the gradients 5 million examples (**inefficient**)
- In Stochastic Gradient Descent (SGD), we consider just **one example** at a time to take a single step.
- Cost will fluctuate over the training examples
- N **steps/iterations** of gradient descent in one epoch

Mini Batch Gradient Descent

- We use a batch of a fixed number of training examples which is less than the actual dataset and call it a mini-batch
- Average cost over the epochs in mini-batch gradient descent fluctuates because we are averaging a small number of examples at a time.
- $N/\text{batch_size}$ **steps/iterations** of gradient descent in one epoch

Batch vs Stochastic vs Mini-Batch



Batch GD

Stochastic GD

Mini-Batch GD

We can divide the dataset of 2000 examples into batches of 500 then it will take ? iterations to complete 1 epoch.

A SIMPLE REGRESSION PROBLEM (NUMPY IMPLEMENTATION)

Dataset Generation

```
1  import numpy as np
2  import matplotlib.pyplot as plt
3
4  #data generation
5  true_w=2
6  true_b=1
7  N=100
8
9  np.random.seed(100)
10 #get N uniformly distributed values
11 x=np.random.rand(N,1)
12 #get N noise values from standard normal distribution
13 epsilon=0.1*np.random.randn(N,1)
14 y=true_w*x+true_b+epsilon
```

Splitting Dataset into train/validation sets

```
16 #Splitting Data into train and validation
17 idx=np.arange(N)
18 np.random.shuffle(idx)
19 idx_train=idx[:int(0.8*N)]
20 idx_test=idx[int(0.8*N):]
21 x_train, y_train = x[idx_train],y[idx_train]
22 x_val, y_val = x[idx_test],y[idx_test]
```


WATCH

```
> x_val.shape: (20, 1)
> y_val.shape: (20, 1)
> x_train.shape: (80, 1)
> x_train.shape: (80, 1)
```

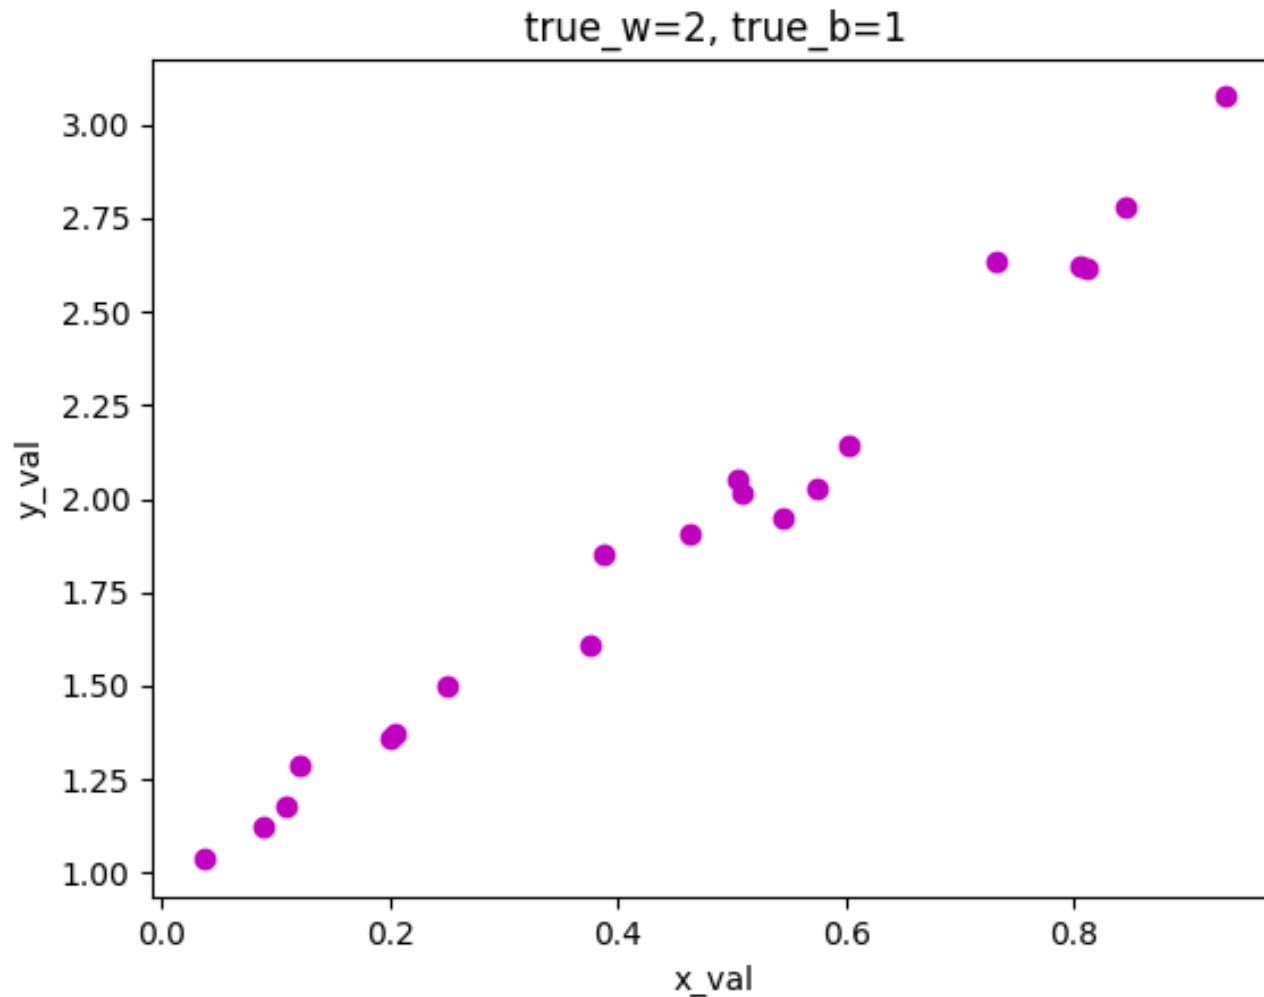
Splitting Dataset into train/validation sets

```
25 #plotting tain and val data
26 plt.figure('1')
27 plt.scatter(x_train,y_train)
28 plt.xlabel('x_train')
29 plt.ylabel('y_train')
30 plt.title(f'true_w={true_w}, true_b={true_b}')
31 plt.figure('2')
32 plt.scatter(x_val,y_val,color = 'm')
33 plt.xlabel('x_val')
34 plt.ylabel('y_val')
35 plt.title(f'true_w={true_w}, true_b={true_b}')
36 plt.show(block=True)
```

Visualizing Datasets



Visualizing Datasets



Training Loop

```
39 #training loop
40 #initializing parameters
41 trainLosses=[]
42 valLosses=[]
43 lr=0.1
44 w=np.random.randn(1)
45 b=np.random.randn(1)
46 for i in range(100):
47     #forward pass
48     yhat=w*x_train+b #note vectorized operation
49     #MSE loss
50     error=yhat-y_train
51     loss= (error**2).mean()
52     trainLosses.append(loss)
53     #computing gradients
54     db=2*error.mean()
55     dw=2*(x_train*error).mean()
56     #weight update
57     b=b-lr*db
58     w=w-lr*dw
```

Validation Loss

```
60     #val MSE loss
61     yhatVal=w*x_val+b
62     errorVal=yhatVal-y_val
63     valLoss= (errorVal**2).mean()
64     valLosses.append(valLoss)
65
66     #stopping condition
67     if(valLoss<0.0001):
68         break
69
70     print(f'train loss={loss}, val loss={valLoss}, w={w}, b={b}')
```


Plots (Regression Fitting)

```
72     #training data plot
73     plt.figure('3')
74     plt.cla()
75     plt.scatter(x_train,y_train)
76     plt.scatter(x_train,yhat)
77     plt.title(f'epoch={i}, loss={loss}, w={w}, b={b}')
78     plt.show(block=False)
79     plt.pause(1)
80
81     #validation data plot
82     plt.figure('4')
83     plt.cla()
84     plt.scatter(x_val,y_val)
85     plt.scatter(x_val,yhatVal)
86     plt.title(f'epoch={i}, ValLoss={valLoss}, w={w}, b={b}')
87     plt.show(block=False)
88     plt.pause(1)
```


Plots (Loss vs Epoch)

```
90     #trainLoss vs Epoch
91     plt.figure('5')
92     plt.cla()
93     plt.plot(trainLosses)
94     plt.xlabel('Epoch')
95     plt.ylabel('trainLoss')
96     plt.title(f'Training Loss Vs Epoch')
97     plt.show(block=False)
98     plt.pause(1)
99
100    #validationLoss vs Epoch
101    plt.figure('6')
102    plt.cla()
103    plt.plot(vallosses,color='m')
104    plt.xlabel('Epoch')
105    plt.ylabel('valLoss')
106    plt.title(f'Validation Loss Vs Epoch')
107    plt.show(block=False)
108    plt.pause(1)
```


Plots (Convergence of regression)




Plots (Convergence of regression)




Plots (Convergence of regression)



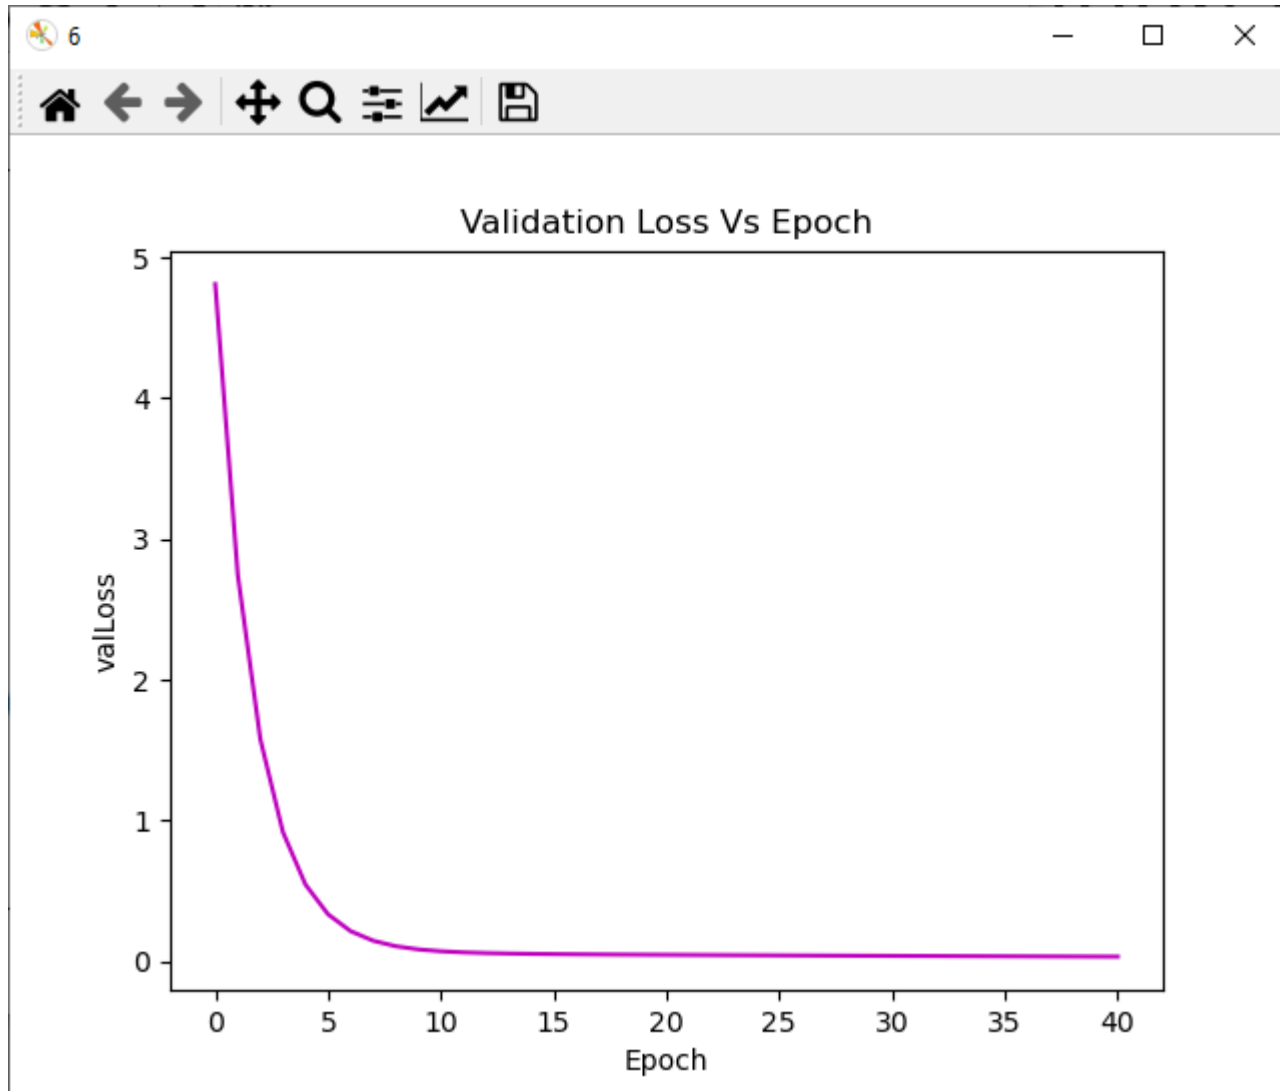
Plots (Performance on Validation Set)



Epoch vs Train Loss



Epoch vs Validation Loss



NN Summary

- Data Set, Training, Validation, Test
- Cost/Loss Function
 - MSE Loss for regression
- Training Loop
- Optimizer
- Parameters
- Learning Rate
- Epoch
- Batch
- Loading/Saving Model

Home Task

- Compare learning curves for different values of learning rate
- Convert code from Batch Gradient Descent to Stochastic Gradient Descent and compare learning curves